

Specifying and Verifying Systems in TLA⁺

Stephan Merz

<http://www.loria.fr/~merz/>

INRIA Nancy & LORIA
Nancy, France

INSTITUT NATIONAL
DE RECHERCHE
EN INFORMATIQUE
ET EN AUTOMATIQUE




INRIA

centre de recherche NANCY - GRAND-EST



ICTAC School on Software Engineering
Natal, Brazil, August 2010




PhD 1972 (Brandeis University), Mathematics


- Mitre Corporation, 1962–65
- Marlboro College, 1965–69
- Massachusetts Computer Associates, 1970–77
- SRI International, 1977–85
- Digital Equipment Corporation/Compaq, 1985–2001
- Microsoft Research, since 2001

Pioneer of distributed algorithms

- Natl. Academy of Engineering, PODC Influential Paper Award, ACM SIGOPS Hall of Fame, LICS Award, IEEE John v. Neumann medal, ...
- honorary doctorates (Rennes, Kiel, Lausanne, Lugano, Nancy)



- formal language for describing and reasoning about distributed and concurrent systems
- based on mathematical logic and set theory plus linear time temporal logic TLA
- book: Addison-Wesley, 2003 (free download for personal use)
- supported by tool set (TLA⁺ toolbox)



- formal language for describing and reasoning about distributed and concurrent systems
- based on mathematical logic and set theory plus linear time temporal logic TLA
- book: Addison-Wesley, 2003 (free download for personal use)
- supported by tool set (TLA⁺ toolbox)

Some other publications

- Y. Yu, P. Manolios, L. Lamport: *Model checking TLA⁺ Specifications*. CHARME 1999, pp. 54-66, LNCS 1703.
- S. Merz: *The Specification Language TLA⁺*. In: *Logics of Specification Languages* (D. Bjørner, M. Henson, eds.), Springer 2008, pp. 401-451.
- K. Chaudhuri, D. Doligez, L. Lamport, S. Merz: *Verifying Safety Properties with the TLA⁺ Proof System*. IJCAR 2010, pp. 142-148, LNCS 6173

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA^+
- 4 Verification of TLA^+ models
- 5 Structuring models: refinement and (de)composition

Why formal specifications?

- describe, analyze, and reason about systems
 - algorithms, protocols, embedded systems, ...
- at different levels of abstraction
 - requirements, high-level design, component design, code
- verify models throughout development
 - gradually introduce design decisions, detect errors early
- focus on different aspects of system under development
 - functional correctness, performance, resource usage, security, ...
- use different analysis and verification techniques
 - simulation, model checking, static analysis, theorem proving, ...

Classes of formal specification languages

- intended for different kinds of systems and properties
 - ▶ sequential algorithms
 - ▶ interactive, reactive, distributed systems
 - ▶ real-time and hybrid systems
 - ▶ security-sensitive protocols and systems
- based on different formal foundations
 - ▶ logics: first-order, higher-order, temporal, ...
 - ▶ automata: transition systems, timed, hybrid automata, ...
 - ▶ abstract machines: λ -calculus, rewriting, process calculi, ...
- based on different specification styles
 - ▶ property-oriented: list desired correctness properties
 - ▶ model-based: describe abstract realization of intended system
- TLA⁺ : reactive systems, temporal logic, model-based

Simple TLA⁺ specification: the hour clock

```

MODULE HourClock
EXTENDS Naturals
VARIABLE hr

HCini     $\triangleq$   $hr \in 0..23$ 
HCnxt     $\triangleq$   $hr' = \text{IF } hr = 23 \text{ THEN } 0 \text{ ELSE } hr + 1$ 
HCsafe     $\triangleq$   $HCini \wedge \Box [HCnxt]_{hr}$ 

THEOREM  $HCsafe \Rightarrow \Box HCini$ 
```


HourClock module: some observations

- Structure of the module

- ▶ list of declarations, definitions, and assertions
- ▶ *hr* a state variable
- ▶ *HCini* a state predicate
- ▶ *HCnxt* a transition predicate (an action)
- ▶ *HCsafe* a temporal formula

- Formula *HCsafe* holds of a behavior if

- ▶ initial state satisfies *HCini*
- ▶ every transition satisfies *HCnxt* – or leaves *hr* unchanged
- ▶ clock may stop ticking, but may not fail in any other way

- Module asserts a theorem

- ▶ any run of the hour clock always satisfies *HCini*
- ▶ theorem can be verified using TLC, the TLA⁺ model checker

Non-stopping hour clock

- Formulas $Init \wedge \Box[Next]_v$ describe a state machine

- ▶ initial state of system
- ▶ allowed state transitions, including “stuttering”

- Fairness conditions rule out infinite stuttering

- ▶ fairness: action must occur if it is possible “often enough”
- ▶ an hour clock that never stops is specified by

$$HC \triangleq HC_{safe} \wedge WF_{hr}(HC_{nxt})$$

- ▶ TLC verifies the theorem

$$HC \Rightarrow \forall k \in 0..23 : \Box\Diamond(hr = k)$$

- Standard form of TLA⁺ specifications: $Init \wedge \Box[Next]_v \wedge L$

- ▶ $Init$ system’s initial condition
- ▶ $Next$ system’s next-state relation (usually a disjunction)
- ▶ L fairness conditions (for disjuncts of $Next$)

Outline

- 1 Motivation and introductory example
- 2 **Transition systems and their properties**
 - Transition systems and their runs
 - Properties of runs
- 3 System Specification in TLA⁺
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

Transformational vs. reactive systems

- Transformational systems (sequential algorithms)
 - ▶ compute result from initially given input
 - ▶ semantics: relation between states before and after computation
 - ▶ partial correctness + termination
 - ▶ computational models: Turing machines, λ -calculus, ...
- Reactive systems (operating systems, controllers, ...)
 - ▶ repeated interaction between environment and system
 - ▶ semantics: (possibly) infinite sequences of states
 - ▶ safety properties: **something bad never happens**
 - ▶ liveness properties: **something good happens eventually**
 - ▶ computational models: transition systems

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
 - Transition systems and their runs
 - Properties of runs
- 3 System Specification in TLA⁺
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

Transition systems

Definition

A transition system $\mathcal{T} = (Q, I, \delta)$ is given by

- a (finite or infinite) set Q of **states**,
- a set $I \subseteq Q$ of **initial states**,
- a total **transition relation** $\delta \subseteq Q \times Q$.

An **execution** of \mathcal{T} is an infinite sequence $\rho = q_0 q_1 q_2 \dots$ where $(q_i, q_{i+1}) \in \delta$ holds for all $i \in \mathbb{N}$.

A **run** of \mathcal{T} is an execution that starts in an initial state of \mathcal{T} .

A state $q \in Q$ is **reachable** if it appears in some run of \mathcal{T} .

- similar to non-deterministic finite automaton; no final states
- totality of δ for technical convenience: restrict to infinite runs
- easy to ensure if δ is reflexive: **stuttering closure**
- deadlocks must be modeled explicitly

Variation: labeled transition systems

- Distinguish several kinds of transitions
 - ▶ transitions of system and environment
 - ▶ transitions of different processes
 - ▶ time elapse vs. system move

Definition

In a labeled transition system $\mathcal{T} = (\mathcal{A}, Q, I, \delta)$:


- \mathcal{A} is a set of labels (actions),
- $\delta = (\delta_A)_{A \in \mathcal{A}}$ is a family of relations $\delta_A \subseteq Q \times Q$ indexed by \mathcal{A} .

Action A is **enabled** at state q if $(q, q') \in \delta_A$ for some q' .

Action A is **taken** at position i in run $\rho = q_0 q_1 q_2 \dots$ if $(q_i, q_{i+1}) \in \delta_A$.

- Remarks
 - ▶ at every state, some action should be enabled
 - ▶ actions will be used later for defining **fairness properties**

Example



- q_0 and q_4 are initial states
- q_5 is unreachable
- A and B are enabled at q_0
- A is not enabled at q_1


Exercises: transition system representations

1 A “stopwatch” program:

```
var  $x, y$  : integer = 0, 0;  
cobegin  
   $\alpha$  : while  $y = 0$  do  $\beta$  :  $x := x + 1$  end  ||   $\gamma$  :  $y := 1$   
coend
```

Describe informally the runs of the transition system.

2 A toy railway:



Model the behavior of the trains and the signals as a transition system.

Fairness conditions (informally)

- Transition systems are non-deterministic
 - ▶ choice between transitions of different processes
 - ▶ non-deterministic choice resulting from abstraction in the model
e.g.: abstraction of concrete train positions by sections
 - ▶ non-determinism generates “unfair” executions
e.g.: never execute action γ of stopwatch program
 - ▶ rule out such executions by global constraints
- Fairness conditions exclude unfair executions
 - ▶ action should be eventually taken if **often enough** enabled
 - ▶ **weak fairness**: continuously enabled
 - ▶ **strong fairness**: infinitely often enabled
- **Note**: fairness does not restrict local non-determinism

Fairness conditions (formally)

Definition (weak and strong fairness)

Let $\rho = q_0q_1q_2 \dots$ be an execution of a labeled transition system and A be an action.

- ① ρ is **weakly fair** w.r.t. A if for all $m \in \mathbb{N}$, if A is enabled at all states q_n for $n \geq m$ then A is taken at some position $n \geq m$.
- ② ρ is **strongly fair** w.r.t. A if for all $m \in \mathbb{N}$, if A is enabled at infinitely many states q_n for $n \geq m$ then A is taken at some position $n \geq m$.

● Fairness conditions rule out executions

- ▶ ρ is not weakly fair w.r.t. A if there exists some $m \in \mathbb{N}$ such that A is forever enabled from position m onward but never taken
- ▶ ρ is not strongly fair w.r.t. A if A is infinitely often enabled but taken only finitely often


Fairness: examples and exercises

- Are the following runs fair w.r.t. A and B ?

- $\rho_1 = q_0 q_1 q_0 q_1 \dots$
- $\rho_2 = q_0 q_2 q_3 q_2 q_3 \dots$
- $\rho_2 = q_4 q_2 q_4 q_2 \dots$

- Exercises

- ρ is weakly fair w.r.t. A iff either A is taken infinitely often in ρ or A is infinitely often disabled.
- ρ is strongly fair w.r.t. A iff either A is taken infinitely often in ρ or A is only finitely often enabled.
- If ρ is strongly fair w.r.t. A then ρ is weakly fair w.r.t. A .
- If ρ is (strongly/weakly) fair w.r.t. A then so is any suffix of ρ .



Outline

- 1 Motivation and introductory example
- 2 **Transition systems and their properties**
 - Transition systems and their runs
 - **Properties of runs**
- 3 System Specification in TLA^+
- 4 Verification of TLA^+ models
- 5 Structuring models: refinement and (de)composition

Properties of transition systems

- Properties evaluated over runs

- ▶ two trains are never simultaneously in section onbridge
- ▶ no train waits forever at the signal
- ▶ if both trains wait at signal then each will enter the bridge at most once before the other

- Properties related to the transition structure

- ▶ some initial state is reachable from all states
- ▶ actions A and B are in conflict, resp. are independent
- ▶ two processes can cooperate to starve a third process

- Here: consider first class of properties

- ▶ define properties as sets of runs
- ▶ system correctness: $Runs(\mathcal{T}) \subseteq \Phi$

Safety and liveness properties (informally)

- Safety properties: “something bad never happens”
 - ▶ trains are never simultaneously on the bridge
 - ▶ data arrives at receiver in FIFO order
- Liveness properties: “something good happens eventually”
 - ▶ train waiting at signal will eventually enter onbridge
 - ▶ every data item sent will eventually be received
 - ▶ process *A* executes infinitely often
- Two fundamental classes of properties of runs
 - ▶ generalization of partial correctness and termination
 - ▶ the following property is neither (pure) safety nor (pure) liveness:
*the train remains at the signal until it enters onbridge,
which will happen eventually*

Formal definitions (Alpern & Schneider 1985)

Definition

Φ is a **safety property** if for every infinite state sequence σ :

$$\sigma \in \Phi \quad \text{iff} \quad \text{for all } n \in \mathbb{N} \text{ there is } \tau \text{ s.t. } \sigma[..n] \circ \tau \in \Phi.$$

- σ violates a safety property if there exists a finite prefix $\sigma[..n]$ that cannot be extended to an infinite sequence satisfying Φ
- “bad things” are observable after some finite time

Definition

Φ is a **liveness property** if for every finite state sequence ρ there exists τ s.t. $\rho \circ \tau \in \Phi$.

- liveness properties do not constrain finite prefixes
- the “good thing” may still occur later


Safety/liveness: examples and exercises

- 1 The set of runs of a transition system is a safety property.
- 2 Weak or strong fairness conditions are liveness properties.
- 3 If all Φ_i ($i \in I$) are safety properties then $\bigcap_{i \in I} \Phi_i$ is a safety property.
- 4 If Φ is a liveness property and $\Phi \subseteq \Psi$ then Ψ is a liveness property.
- 5 The trivial property containing all state sequences is the only property that is both a safety and a liveness property.
- 6 Given property Φ , its **safety closure**
$$\mathcal{C}(\Phi) = \{ \sigma : \text{for all } n \in \mathbb{N} \text{ there is } \tau \text{ s.t. } \sigma[..n] \circ \tau \in \Phi \}$$
is the smallest safety property containing Φ .
- 7 For any property Φ there exist a safety property S_Φ and a liveness property L_Φ such that $\Phi = S_\Phi \cap L_\Phi$.

Machine closure

- Specifications: transition system plus liveness properties

- ▶ liveness properties can impose additional constraints on transitions
- ▶ example: liveness property that requires infinitely many visits of s_0



Definition

Let S be a safety property and L be any property.

The pair (S, L) is **machine closed** if $\mathcal{C}(S \cap L) = S$.

- ▶ **Idea:** L constrains infinite, but not finite runs satisfying S
- ▶ **Show:** if (S, L) is machine closed and Φ is a safety property then $S \cap L \subseteq \Phi$ iff $S \subseteq \Phi$
- ▶ **Theorem:** specifications given by a transition system and a countable number of fairness condition are machine closed

- Transition systems: semantic basis for reactive systems
 - ▶ transition relation describes possible steps
 - ▶ fairness conditions: global constraints on non-deterministic choices
- Properties of runs: sets of state sequences
 - ▶ “linear-time” correctness properties, ignore branching properties
 - ▶ dichotomy of safety and liveness properties

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺**
 - Temporal Logic of Actions
 - Set theory: specifying data structures in TLA⁺
 - The module language of TLA⁺
 - Specification styles: modeling a queue
 - Case study: a resource allocator (part 1)
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

The Specification Language TLA⁺

- Temporal Logic of Actions (TLA)
 - ▶ linear-time temporal logic
 - ▶ formulas represent systems and properties
 - ▶ **stuttering invariance**: basis for refinement and composition
- Zermelo-Fränkel set theory: classical mathematics
 - ▶ specification of data structures in set theory
 - ▶ high levels of abstraction and expressiveness
 - ▶ explicit definitions of data and operations (e.g., communication)
- Module language
 - ▶ structuring of specifications: composition and hiding
 - ▶ standard library of common structures, user extensible

Temporal logics in computer science

- Originally: temporal relations in natural language
 - ▶ *Yesterday she said that she'd come tomorrow, so she should come today.*
 - ▶ A. Prior: *Past, present, and future*. Oxford University Press, 1967
- 1977: express properties of reactive systems
 - ▶ A. Pnueli: *The temporal logic of programs*. FOCS'77
- 1981: automatic verification of finite-state systems
 - ▶ E.M. Clarke, E.A. Emerson: *Synthesis of synchronization skeletons for branching time temporal logic*. Logics of Programs, 1981
 - ▶ J.P. Queille, J. Sifakis: *Specification and verification of concurrent systems in Cesar*. Intl. Symp. Programming, LNCS 137

System	satisfies	property
	formalized as	
Transition system	is model of	temporal formula

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺**
 - Temporal Logic of Actions
 - Set theory: specifying data structures in TLA⁺
 - The module language of TLA⁺
 - Specification styles: modeling a queue
 - Case study: a resource allocator (part 1)
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

Anatomy of TLA

- Variables and constants

- ▶ (state) variables $v, hr, queue$ represent system state
- ▶ constants $x, k, nProcs$ represent system parameters

- Action formulas: predicates over (pairs of) states

- ▶ represent initial condition, invariants, system transitions

$$hr \in 0..23$$

$$hr' = \text{IF } hr = 23 \text{ THEN } 0 \text{ ELSE } hr + 1$$

$$\exists p \in 1..nProcs : Request(p)$$

- Temporal formulas: predicates over state sequences

- ▶ represent system specifications or properties

$$Init \wedge \Box [Next]_{vars} \wedge WF_{vars}(Exit_1 \vee Exit_2)$$

$$\forall p \in 1..nProcs : request[p] \leadsto grant[p]$$

Syntax of action formulas

- First-order formulas (over given signature) containing

- ▶ constants $x, y, k, \dots \in \mathcal{X}$
- ▶ state variables $v, w, \dots \in \mathcal{V}$
- ▶ primed state variables $v', w', \dots \in \mathcal{V}'$

- Examples

- ▶ $v = 0, \quad v > 0 \wedge v' = v - 1, \quad \exists k : v - w' < x + 3 * k$
- ▶ free and bound variables, substitution: as in first-order logic

- State and transition formulas

- ▶ terms: **transition functions**
- ▶ formulas: **transition predicates**
- ▶ formulas without free primed flexible variables: **state formulas**
- ▶ formulas without flexible variables: **constant formulas**

Semantics of action formulas

- First-order interpretation \mathcal{I} for underlying signature
 - ▶ provides a non-empty universe $|\mathcal{I}|$ of values
 - ▶ interprets function and predicate symbols: $0, +, <, \in, \dots$
- Interpretations of constants and variables
 - ▶ valuations of state variables: $\text{states } s : \mathcal{V} \rightarrow |\mathcal{I}|$
 - ▶ valuation of constants $\xi : \mathcal{X} \rightarrow |\mathcal{I}|$
- Interpretation of action formulas $\llbracket A \rrbracket_{s,t}^{\xi} \in \{\text{tt}, \text{ff}\}$
 - ▶ standard inductive definition
 - ▶ s, t interpret unprimed and primed flexible variables
 - ▶ ξ interprets rigid variables
- Interpretation of state formulas does not depend on second state

Notations for action formulas (1)

- “Priming” of state formulas e

- e' action formula obtained by replacing every variable v by v'
- rename bound variables as necessary to avoid confusion

$$\begin{aligned}(v + 1)' &\equiv v' + 1 \\ (\exists k : v = k + w)' &\equiv \exists k : v' = k + w' \\ (\exists v' : v = v' + w)' &\equiv \exists vp : v' = vp + w'\end{aligned}$$

- UNCHANGED $e \triangleq e' = e$

- Bracketed action formulas

$$\begin{aligned}[A]_e &\triangleq A \vee e' = e && \text{A or stuttering (w.r.t. } e) \\ \langle A \rangle_e &\triangleq A \wedge \neg(e' = e) && \text{A and change of } e\end{aligned}$$

- Note: $\langle A \rangle_e \equiv \neg[\neg A]_e$ $\neg\langle A \rangle_e \equiv [\neg A]_e$
 $[A]_e \equiv \neg\langle\neg A\rangle_e$ $\neg[A]_e \equiv \langle\neg A\rangle_e$

Notations for action formulas (2)

- Enabledness of actions

- $\text{ENABLED } A \triangleq \exists v'_1, \dots, v'_n : A$
(v'_1, \dots, v'_n all free primed state variables in A)
- $\llbracket \text{ENABLED } A \rrbracket_s^{\xi} = \text{tt}$ iff $\llbracket A \rrbracket_{s,t}^{\xi} = \text{tt}$ for some state t
- will be used to define fairness conditions in TLA

- Sequential composition of actions (atomic)

- $A \cdot B \equiv \exists v''_1, \dots, v''_n : A[v''_1/v'_1, \dots, v''_n/v'_n] \wedge$
 $B[v''_1/v_1, \dots, v''_n/v_n]$
(v_1, \dots, v_n all free state variables in A or B)
- $\llbracket A \cdot B \rrbracket_{s,t}^{\xi} = \text{tt}$ iff $\llbracket A \rrbracket_{s,u}^{\xi} = \text{tt}$ and $\llbracket B \rrbracket_{u,t}^{\xi} = \text{tt}$ for some state u
- sequential composition of A and B as single atomic action

Temporal formulas of TLA (to be completed)

- interpreted over ω -sequence $\sigma = s_0s_1 \dots$ of states and valuation ξ

- ▶ $\sigma, \xi \models F$ if F holds over (σ, ξ)
- ▶ write $\sigma[n..]$ for suffix $s_ns_{n+1} \dots$

- Inductive definition of temporal formulas

state formulas $\sigma, \xi \models P$ iff $\llbracket P \rrbracket_{s_0}^{\xi} = \text{tt}$

Boolean combinations $\sigma, \xi \models F \wedge G$ iff $\sigma, \xi \models F$ and $\sigma, \xi \models G$, etc.

first-order quantifiers $\sigma, \xi \models \exists x : F$ iff $\sigma, \xi \models F$ for some $\xi \sim_x \xi$

always $\sigma, \xi \models \Box F$ iff $\sigma[n..], \xi \models F$ for all $n \in \mathbb{N}$

always square A sub e $\sigma, \xi \models \Box[A]_e$ iff $\llbracket [A]_e \rrbracket_{s_n, s_{n+1}}^{\xi} = \text{tt}$ for all $n \in \mathbb{N}$

- ▶ **Note:** $[A]_e$ and $\Box A$ are not temporal formulas (for action A)
- ▶ if ξ is unimportant, just write $\sigma \models F$

Notations for temporal formulas

- “eventually” operator $\Diamond F \triangleq \neg \Box \neg F$
 - ▶ semantics: $\sigma \models \Diamond F$ iff $\sigma[n..] \models F$ for some $n \in \mathbb{N}$
- “eventually angle A sub e ” $\Diamond \langle A \rangle_e \triangleq \neg \Box [\neg A]_e$
 - ▶ some future state transition satisfies $\langle A \rangle_e$
 - ▶ again, $\langle A \rangle_e$ and $\Diamond A$ are not temporal formulas
- “leadsto” operator $F \leadsto G \equiv \Box (F \Rightarrow \Diamond G)$
 - ▶ every occurrence of F is followed by some occurrence of G

Example: semantics of temporal formulas

Which of the following formulas hold in this behavior?

										→	
x	0	0	3	7	0	1	1	0	2	...	(always $\neq 0$)
y	1	1	0	0	0	0	3	4	0	...	(always $= 0$)

- $\Box \neg (x = 0 \wedge y = 0)$
- $\Box [x = 0 \Rightarrow y' = 0]_{x,y}$
- $\Diamond (x = 7 \wedge y = 0)$
- $\Diamond \langle y = 0 \wedge x' = 0 \rangle_y$
- $\Box \Diamond (y \neq 0)$
- $\Diamond \Box (x = 0 \Rightarrow y \neq 0)$
- $\Diamond \Box [\text{FALSE}]_y$

Infinitely often and eventually always

- $\Box\Diamond F$ asserts that F holds infinitely often

- ▶ $\sigma \models \Box\Diamond F$ iff for all $m \in \mathbb{N}$ there is $n \geq m$ such that $\sigma[n..] \models F$
- ▶ similarly: $\Box\Diamond\langle A \rangle_e$ asserts that $\langle A \rangle_e$ occurs infinitely often

- $\Diamond\Box F$ asserts that F continues to hold from some time on

- ▶ equivalently, F is false only finitely often
- ▶ similarly: $\Diamond\Box[A]_e$ asserts that eventually only $[A]_e$ actions occur

- Equivalences

$$\neg\Box\Diamond F \equiv \Diamond\Box\neg F$$

$$\neg\Diamond\Box F \equiv \Box\Diamond\neg F$$

$$\Diamond\Box\Diamond F \equiv \Box\Diamond F$$

$$\Box\Diamond\Box F \equiv \Diamond\Box F$$

Fairness of actions in TLA

- Weak fairness

- ▶ action will be taken eventually if it is continuously enabled

$$WF_e(A) \triangleq \Box(\Box \text{ENABLED } \langle A \rangle_e \Rightarrow \Diamond \langle A \rangle_e)$$

- Strong fairness

- ▶ action will be taken eventually if it is infinitely often enabled

$$SF_e(A) \triangleq \Box(\Box \Diamond \text{ENABLED } \langle A \rangle_e \Rightarrow \Diamond \langle A \rangle_e)$$

- Equivalent conditions

$$WF_e(A) \equiv \Diamond \Box \text{ENABLED } \langle A \rangle_e \Rightarrow \Box \Diamond \langle A \rangle_e$$

$$WF_e(A) \equiv \Box \Diamond \neg \text{ENABLED } \langle A \rangle_e \vee \Box \Diamond \langle A \rangle_e$$

$$SF_e(A) \equiv \Box \Diamond \text{ENABLED } \langle A \rangle_e \Rightarrow \Box \Diamond \langle A \rangle_e$$

$$SF_e(A) \equiv \Diamond \Box \neg \text{ENABLED } \langle A \rangle_e \vee \Box \Diamond \langle A \rangle_e$$

Example: “stopwatch” program in TLA⁺

```

MODULE Stopwatch
EXTENDS Naturals
VARIABLES pc1, pc2, x, y

Init ≜ pc1 = "alpha" ∧ pc2 = "gamma" ∧ x = 0 ∧ y = 0
A    ≜ ∧ pc1 = "alpha"
      ∧ pc'1 = IF y = 0 THEN "beta" ELSE "stop"
      ∧ UNCHANGED ⟨pc2, x, y⟩
B    ≜ ∧ pc1 = "beta" ∧ pc'1 = "alpha"
      ∧ x' = x + 1 ∧ UNCHANGED ⟨pc2, y⟩
G    ≜ ∧ pc2 = "gamma" ∧ pc'2 = "stop"
      ∧ y' = 1 ∧ UNCHANGED ⟨pc1, x⟩
vars ≜ ⟨pc1, pc2, x, y⟩
Spec ≜ Init ∧ □[A ∨ B ∨ G]vars ∧ WFvars(A ∨ B) ∧ WFvars(G)

```

```

var x, y : integer = 0, 0;
cobegin
  α : while y = 0 do
    β : x := x + 1
  end
||
  γ : y := 1
coend

```

- explicit encoding of control structure
- process structure not obvious from TLA⁺ representation

Aside: PlusCal (Lamport, 2006)

```
-- algorithm Peterson {  
variables req  = [id ∈ Node ↦ FALSE],  
          turn = 0  
process (Proc ∈ Node) {  
  ncs: while (TRUE) {  
    skip;  
    rq:   req[self] := TRUE; turn := other(self);  
    try:  await (turn = self ∨ ¬lock[other(self)]);  
    cs:   skip;  
    lv:   req[self] := FALSE  
  } } }
```

- High-level language for multi-process algorithms
 - ▶ atomicity indicated through labels
 - ▶ algorithm embedded in TLA⁺ module
- Translation to TLA⁺: simulation and model checking

Stuttering invariance

- Action formulas must be “protected” in TLA: $\Box[A]_t, \Diamond\langle A \rangle_t$
 - ▶ finite repetition of identical states (up to t) preserves truth
 - ▶ this observation extends to all TLA formulas
- Stuttering equivalence (\approx)
 - ▶ $\Downarrow(s_0s_1\ldots) \triangleq$ IF $\forall k \in \mathbb{N} : s_k = s_0$
THEN $s_0s_1\ldots$
ELSE LET $m = \min\{k \in \mathbb{N} : s_k \neq s_0\}$
IN $s_0 \circ \Downarrow(s_ms_{m+1}\ldots)$
 - ▶ $\sigma \approx \tau$ iff $\Downarrow\sigma = \Downarrow\tau$

Stuttering invariance

- Action formulas must be “protected” in TLA: $\Box[A]_t, \Diamond\langle A\rangle_t$
 - ▶ finite repetition of identical states (up to t) preserves truth
 - ▶ this observation extends to all TLA formulas
- Stuttering equivalence (\approx)
 - ▶ $\Downarrow(s_0s_1\ldots) \triangleq$ IF $\forall k \in \mathbb{N} : s_k = s_0$
THEN $s_0s_1\ldots$
ELSE LET $m = \min\{k \in \mathbb{N} : s_k \neq s_0\}$
IN $s_0 \circ \Downarrow(s_ms_{m+1}\ldots)$
 - ▶ $\sigma \approx \tau$ iff $\Downarrow\sigma = \Downarrow\tau$

Theorem (stuttering invariance)

For any temporal formula F and behaviors $\sigma \approx \tau$:

$$\sigma, \xi \models F \quad \text{iff} \quad \tau, \xi \models F$$

- Fundamental for refinement and composition

Example: hour-minute clock

- Extension of hour clock by a minute hand

```

————— MODULE HourMinuteClock —————
EXTENDS Naturals, HourClock
VARIABLE min

HMCini     $\triangleq$   HCini  $\wedge$  min  $\in$  0..59
Min       $\triangleq$   min' = IF min = 59 THEN 0 ELSE min + 1
Hr        $\triangleq$   (min = 59  $\wedge$  HCnxt)  $\vee$  (min < 59  $\wedge$  hr' = hr)
HMCnxt     $\triangleq$   Min  $\wedge$  Hr
HMC       $\triangleq$   HMCini  $\wedge$   $\Box$ [HMCnxt] $\langle$ hr,min $\rangle$   $\wedge$  WF $\langle$ hr,min $\rangle$ (HMCnxt)

THEOREM  HMC  $\Rightarrow$  HC

```

- ▶ stuttering invariance is essential for refinement as implication
- ▶ no formal difference between specifications and properties

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺**
 - Temporal Logic of Actions
 - **Set theory: specifying data structures in TLA⁺**
 - The module language of TLA⁺
 - Specification styles: modeling a queue
 - Case study: a resource allocator (part 1)
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

Logical language underlying TLA⁺

- Specification language TLA⁺: based on ZF set theory

- ▶ standard logical basis for formalizing mathematics
- ▶ high levels of abstraction and expressiveness

e.g., Riemann integral in 15 lines of TLA⁺

- ▶ set-theoretical language, no types

$5 = \{\}$ or $17 \wedge \text{"abc"}$ are well-formed formulas
but (of course!) we don't know if they are true

- Formally: logical basis

- ▶ binary predicate symbol \in
- ▶ choice operator (binder) $\text{CHOOSE } x : P$

- Hilbert's choice operator

- ▶ $\text{CHOOSE } x : P$ denotes an arbitrary value satisfying P
- ▶ or an arbitrary, fixed value if no such value exists

Set-theoretic constructions

- Characteristic axioms of Hilbert's choice

$$(\exists x : P(x)) \equiv P(\text{CHOOSE } x : P(x)) \quad [\text{choice}]$$

$$(\forall x : P \equiv Q) \Rightarrow (\text{CHOOSE } x : P) = (\text{CHOOSE } x : Q) \quad [\text{determinacy}]$$

- Examples: set-theoretical constructions and theorems

- ▶ generalized union

$$\text{UNION } S \stackrel{\Delta}{=} \text{CHOOSE } M : \forall x : x \in M \equiv \exists T \in S : x \in T$$

- ▶ no set contains all elements

$$(\text{CHOOSE } x : x \notin S) \notin S$$

- ▶ Russell's paradoxical set is not well-defined

$$(\text{CHOOSE } S : \forall x : x \in S \equiv x \notin x) = (\text{CHOOSE } x : \text{ff})$$

- Exercises: define the following constructions

- ▶ subset relation $S \subseteq T$

- ▶ set intersection $S \cap T$

- ▶ powerset $\text{SUBSET } S$: set containing all subsets of S

- ▶ set comprehensions $\{x \in S : P(x)\}$ and $\{e(x) : x \in S\}$

Choice vs. non-determinism

- Two formulas for specifying resource allocation

$$Alloc_{nd} \triangleq$$

$$\wedge owner = NoProcess$$

$$\wedge waiting \neq \{\}$$

$$\wedge owner' \in waiting$$

$$\wedge waiting' = waiting \setminus \{owner'\}$$

$$Alloc_{ch} \triangleq$$

$$\wedge owner = NoProcess$$

$$\wedge waiting \neq \{\}$$

$$\wedge owner' = \text{CHOOSE } p : p \in waiting$$

$$\wedge waiting' = waiting \setminus \{owner'\}$$

- enabledness: resource is free, some process is waiting
- effect: assign the resource to some waiting process
- $Alloc_{nd}$ produces $Card(waiting)$ successor states
- $Alloc_{ch}$ produces single successor state for some fixed process

CHOOSE operator is deterministic: “arbitrary, but fixed element”

Functions in TLA⁺

- TLA⁺ introduces functions as primitive values

$[S \rightarrow T]$	set of functions with domain S and codomain T
$\text{DOMAIN } f$	domain of function f
$f[e]$	application of function f to argument e
$[x \in S \mapsto e]$	function with domain S mapping x to e
$[f \text{ EXCEPT } ![t] = e]$	function override: like f , but argument t maps to e $[f \text{ EXCEPT } ![t] = @ + e] = [f \text{ EXCEPT } ![t] = f[t] + e]$

- Characteristic property of functional values

$$f = [x \in \text{DOMAIN } f \mapsto f[x]]$$

- Note

- ▶ value $f[x]$ unspecified if $x \notin \text{DOMAIN } f$
- ▶ notations extend to several arguments: $[x \in S, y \in T \mapsto x + y]$

Recursive functions

- Recursive functions can be defined using choice

$$fact \triangleq \text{CHOOSE } f : f = [n \in \text{Nat} \mapsto \text{IF } n = 0 \text{ THEN } 1 \text{ ELSE } n * f[n - 1]]$$

- Abbreviated notation

$$fact[n \in \text{Nat}] \triangleq \text{IF } n = 0 \text{ THEN } 1 \text{ ELSE } n * fact[n - 1]$$

- Potential pitfalls

- ▶ existence of such a function should be justified
- ▶ no implicit commitment to least fixed point semantics

Numbers in TLA⁺ (1)

- Natural numbers defined using choice from Peano axioms

$$\begin{aligned} \text{PeanoAxioms}(N, Z, Sc) &\triangleq \\ &\wedge Z \in N \\ &\wedge Sc \in [N \rightarrow N] \\ &\wedge \forall n \in N : (n \neq Z) \equiv (\exists m \in N : n = Sc[m]) \\ &\wedge \forall S \in \text{SUBSET } N : Z \in S \wedge (\forall n \in S : Sc[n] \in S) \Rightarrow S = N \\ \text{Succ} &\triangleq \text{CHOOSE } Sc : \exists N, Z : \text{PeanoAxioms}(N, Z, Sc) \\ \text{Nat} &\triangleq \text{DOMAIN Succ} \\ \text{Zero} &\triangleq \text{CHOOSE } Z : \text{PeanoAxioms}(\text{Nat}, Z, \text{Succ}) \end{aligned}$$

- ▶ existence of such values can be justified from set theory

Numbers in TLA⁺ (2)

- Arithmetic operations defined recursively (simplified)

$$\begin{array}{lll} \text{pred}(x) & \triangleq & \text{CHOOSE } y : x = \text{Succ}[y] \\ \text{plus}[x \in \text{Nat}, y \in \text{Nat}] & \triangleq & \text{IF } y = 0 \text{ THEN } x \text{ ELSE } \text{Succ}[\text{plus}[x, \text{pred}(y)]] \\ x + y & \triangleq & \text{plus}[x, y] \\ x \leq y & \triangleq & \exists z : y = x + z \end{array}$$

- Numbers and intervals

$$\begin{array}{l} 0 \triangleq \text{Zero}, \quad 1 \triangleq \text{Succ}[0], \quad 2 \triangleq \text{Succ}[1], \quad \dots \\ i..j \triangleq \{n \in \text{Nat} : i \leq n \wedge n \leq j\} \end{array}$$

- Integer and real numbers

- defined similarly as supersets of *Nat*
- arithmetic operations agree: $3.5 + 2.5 \in \text{Nat}$

Tuples and sequences

- Tuples and sequences are represented as functions

$$\langle e_1, \dots, e_n \rangle \triangleq [i \in 1..n \mapsto \text{IF } i = 1 \text{ THEN } e_1 \dots \text{ELSE } e_n]$$

- Standard operators involving sequences

$$\begin{aligned} \text{Seq}(S) &\triangleq \text{UNION } \{[1..n \rightarrow S] : n \in \text{Nat}\} \\ \text{Len}(s) &\triangleq \text{CHOOSE } n \in \text{Nat} : \text{DOMAIN } s = 1..n \\ \text{Head}(s) &\triangleq s[1] \\ \text{Tail}(s) &\triangleq [i \in 1..(\text{Len}(s) - 1) \mapsto s[i + 1]] \\ s \circ t &\triangleq [i \in 1..(\text{Len}(s) + \text{Len}(t)) \mapsto \\ &\quad \text{IF } i \leq \text{Len}(s) \text{ THEN } s[i] \text{ ELSE } t[i - \text{Len}(s)]] \\ \text{Append}(s, e) &\triangleq s \circ \langle e \rangle \end{aligned}$$

- **Question:** What are $\text{Head}(\langle \rangle)$ and $\text{Tail}(\langle \rangle)$?

Exercises

- 1 Define an operator $IsSorted(s)$ such that for any sequence s of numbers, $IsSorted(s)$ is true iff s is sorted.
- 2 Define a function $sort \in [Seq(Real) \rightarrow Seq(Real)]$ such that $sort[s]$ is a sorted sequence containing the same elements as s .
- 3 Give a recursive definition of the $mergesort$ function.
Does $sort = mergesort$ hold for your definitions? Why (why not)?
- 4 Define operators $IsFiniteSet(S)$ and $card(S)$ such that $IsFiniteSet(S)$ holds iff S is a finite set and that $card(S)$ denotes the cardinality of S if S is finite.

Strings and records

- String representations

- ▶ strings are sequences of characters
- ▶ standard operations on sequences apply to strings:
 $\text{"th"} \circ \text{"is"} = \text{"this"}$

- Records: functions whose domain is a finite set of strings

short notation

instead of

$account.bal$

$account["bal"]$

$[account \text{ EXCEPT } !.bal = @ + sum]$

$[account \text{ EXCEPT } !["bal"] = @ + sum]$

$[num \mapsto 1234567, bal \mapsto -321.45]$

$[fld \in \{\text{"num"}, \text{"bal"}\} \mapsto$

IF $fld = \text{"num"}$ THEN 1234567

ELSE - 321.45]

$[x : S, y : T]$

$[fld \in \{\text{"x"}, \text{"y"}\} \rightarrow$

IF $fld = \text{"x"}$ THEN S ELSE $T]$

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺**
 - Temporal Logic of Actions
 - Set theory: specifying data structures in TLA⁺
 - The module language of TLA⁺**
 - Specification styles: modeling a queue
 - Case study: a resource allocator (part 1)
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

TLA⁺ modules: basic ideas

- A TLA⁺ module essentially contains a sequence of
 - ▶ **declarations** of constant and variable parameters
 - ▶ **definitions** of operators (and functions)
 - ▶ **assertions** of assumptions and theorems
- Module hierarchy helps structuring specifications
 - ▶ scoping of parameter and operator names
 - ▶ module extension: import all parameters and definitions
 - ▶ module instantiation: import definitions, must provide parameters
 - ▶ local modules: hide auxiliary definitions
- Meaning of modules
 - ▶ accumulate assumptions: hypotheses for theorems
 - ▶ operators: replace defined symbols by definition body

Principle of unique names

- Identifiers active in a scope cannot be reused
 - ▶ redeclaration or redefinition is illegal
 - ▶ active identifiers must not be used as bound variables

```
MODULE IllegalModule
EXTENDS Naturals
CONSTANTS x, y
|
| m + n       $\triangleq$  ... (* + defined in module Naturals *)
| Foo(y, z)   $\triangleq$   $\exists x : \dots$  (* x and y active in current scope *)
| Nat         $\triangleq$  LET y  $\triangleq$  ... IN ... (* clashes of Nat and y *)
```

- Import of same module via different paths is allowed
- Definitions can be protected from export: LOCAL keyword

Module extension

```
MODULE Foo  
EXTENDS Bar, Baz  
CONSTANTS Data, Compare(_)  
...  
ENDMODULE
```

- Module *Foo* exports

- ▶ symbols declared or defined in module *Foo*
- ▶ (non-local) symbols exported by modules *Bar, Baz*

- Within module *Foo*

- ▶ may use, but not redefine or redeclare symbols exported by *Bar, Baz*

- Equivalent to textual inclusion of modules *Bar, Baz*

Module instantiation: import with renaming

MODULE *Component*

$InChan \triangleq$ INSTANCE *Channel* WITH *Data* \leftarrow *Message*, *chan* \leftarrow *in*

$Chan(c) \triangleq$ INSTANCE *Channel* WITH *Data* \leftarrow *Message*, *chan* \leftarrow *c*

- Use of operators defined in instantiated module

InChan!Send(d) resp. *Chan(in)!Send(d)*

- Special cases

- ▶ identity renaming can be omitted in instantiation
- ▶ instance name can be omitted if only one copy needed
- ▶ LOCAL instantiation is possible

- Operators defined in instantiating module


- ▶ non-local instances are exported for use within later modules
- ▶ symbols declared or defined in module *Channel* are not exported

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺**
 - Temporal Logic of Actions
 - Set theory: specifying data structures in TLA⁺
 - The module language of TLA⁺
 - Specification styles: modeling a queue**
 - Case study: a resource allocator (part 1)
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

Formally describe a FIFO queue

- Basic system component



- ▶ input elements over input “wire” i
- ▶ output over “wire” o , in same order
- ▶ simplifying assumption: subsequent elements are different
- ▶ also assume unbounded queue capacity

- Explore possible specification styles in TLA⁺

- ▶ specify transition system: represent state components
- ▶ encode synchronization and communication

First attempt: lossy queue

MODULE *LossyQueue*

EXTENDS *Sequences*

VARIABLES i, q, o

$LQInit \triangleq q = \langle \rangle \wedge i = o$

$LQEnq \triangleq q' = Append(q, i') \wedge o' = o$

$LQDeq \triangleq q \neq \langle \rangle \wedge o' = Head(q) \wedge q' = Tail(q) \wedge i' = i$

$LQLive \triangleq WF_{q,o}(SQDeq)$

$LQSpec \triangleq LQInit \wedge \Box [LQEnq \vee LQDeq]_{q,o} \wedge LQLive$

- i and o represent interface, q is internal buffer
- interleaving: enqueue and dequeue cannot happen at once
- fairness for dequeue: every element should eventually be output

First attempt: lossy queue

```

MODULE LossyQueue
EXTENDS Sequences
VARIABLES i, q, o

LQInit  $\triangleq q = \langle \rangle \wedge i = o$ 
LQEnq  $\triangleq q' = \text{Append}(q, i') \wedge o' = o$ 
LQDeq  $\triangleq q \neq \langle \rangle \wedge o' = \text{Head}(q) \wedge q' = \text{Tail}(q) \wedge i' = i$ 
LQLive  $\triangleq \text{WF}_{q,o}(\text{SQDeq})$ 
LQSpec  $\triangleq \text{LQInit} \wedge \Box[\text{LQEnq} \vee \text{LQDeq}]_{q,o} \wedge \text{LQLive}$ 

```

- i and o represent interface, q is internal buffer
- interleaving: enqueue and dequeue cannot happen at once
- fairness for dequeue: every element should eventually be output
- **buffer can enqueue same input several times, or not at all**

Synchronous communication, interleaving

```

MODULE SyncQueue
EXTENDS Sequences
VARIABLES i, q, o

SQInit  $\triangleq q = \langle \rangle \wedge i = o$ 
SQEnq  $\triangleq i' \neq i \wedge q' = \text{Append}(q, i') \wedge o' = o$ 
SQDeq  $\triangleq q \neq \langle \rangle \wedge o' = \text{Head}(q) \wedge q' = \text{Tail}(q) \wedge i' = i$ 
SQLive  $\triangleq \text{WF}_{i,q,o}(\text{SQDeq})$ 
SQSpec  $\triangleq \text{SQInit} \wedge \Box[\text{SQEnq} \vee \text{SQDeq}]_{i,q,o} \wedge \text{SQLive}$ 

```

- change of input wire triggers enqueue: i appears in index
- interleaving style and fairness as before
- standard specification style in TLA⁺
- every run of *SQSpec* also satisfies *LQSpec*

Asynchronous communication, interleaving

MODULE AsyncQueue

EXTENDS Sequences

VARIABLES i, q, o, sig

$$AQInit \triangleq q = \langle \rangle \wedge i = o \wedge sig = 0$$

$$AQEnv \triangleq sig = 0 \wedge sig' = 1 \wedge \text{UNCHANGED } \langle q, o \rangle$$

$$AQEnq \triangleq sig = 1 \wedge sig' = 0 \wedge q' = \text{Append}(q, i') \wedge \text{UNCHANGED } \langle i, o \rangle$$

$$AQDeq \triangleq q \neq \langle \rangle \wedge o' = \text{Head}(q) \wedge q' = \text{Tail}(q) \wedge \text{UNCHANGED } \langle i, sig \rangle$$

$$AQLive \triangleq \text{WF}_{i,q,o,sig}(AQEnq) \wedge \text{WF}_{i,q,o,sig}(AQDeq)$$

$$AQSpec \triangleq AQInit \wedge \Box[AQEnv \vee AQEnq \vee AQDeq]_{i,q,o,sig} \wedge AQLive$$

- decouple sender input and acceptance of input by the queue
- explicit “handshake protocol” for enqueueing values
- fairness condition on $AQEnq$ ensures system reaction
- every run of $AQSpec$ also satisfies $LQSpec$

Synchronous communication, non-interleaving

```
MODULE SyncNIQueue
EXTENDS Sequences
VARIABLES i,q,o

SNQInit  $\triangleq q = \langle \rangle \wedge i = o$ 
d(v)  $\triangleq \text{IF } v' = v \text{ THEN } \langle \rangle \text{ ELSE } \langle v' \rangle$ 
SNQEnq  $\triangleq i' \neq i \wedge q \circ d(i) = d(o) \circ q'$ 
SNQDeq  $\triangleq q \neq \langle \rangle \wedge o' = \text{Head}(q) \wedge q \circ d(i) = d(o) \circ q'$ 
SNQLive  $\triangleq \text{WF}_{q,o}(\text{SNQDeq})$ 
SNQSpec  $\triangleq \wedge \text{SNQInit} \wedge \square[\text{SNQEnq}]_i \wedge \square[\text{SNQDeq}]_o$ 
 $\wedge \square[\text{SNQEnq} \vee \text{SNQDeq}]_q \wedge \text{SNQLive}$ 
```

- enqueue and dequeue may happen simultaneously if q non-empty
- separate next-state relation per (group of) variable(s)
- every run of $SQSpec$ also satisfies $SNQSpec$

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺**
 - Temporal Logic of Actions
 - Set theory: specifying data structures in TLA⁺
 - The module language of TLA⁺
 - Specification styles: modeling a queue
 - Case study: a resource allocator (part 1)
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition

Informal requirements

- A set of clients compete for a (fixed, finite) set of resources.
- Whenever a client holds no resources and has no outstanding requests, he can request a set of resources.
- The allocator may allocate a set of available resources to a client that requested them, possibly without completely satisfying the client's request.
- Clients may return resources they hold at any time.
- A client that received all resources he requested must eventually return them (not necessarily at once).

Objectives:

- Clients have exclusive access to resources they hold.
- Every request is eventually satisfied.

A first solution in TLA⁺ (1/4)

MODULE *SimpleAllocator*

EXTENDS *FiniteSet*

CONSTANTS *Clients, Resources*

ASSUME *IsFiniteSet(Resources)*

VARIABLES

unsat, (* *unsat*[*c*] denotes the outstanding requests of client *c* *)

alloc (* *alloc*[*c*] denotes the resources allocated to client *c* *)

TypeInvariant \triangleq

$\wedge \textit{unsat} \in [\textit{Clients} \rightarrow \text{SUBSET } \textit{Resources}]$

$\wedge \textit{alloc} \in [\textit{Clients} \rightarrow \text{SUBSET } \textit{Resources}]$

available \triangleq (* set of resources free for allocation *)
 $\textit{Resources} \setminus (\text{UNION } \{\textit{alloc}[c] : c \in \textit{Clients}\})$

A first solution in TLA⁺ (2/4)

$Init \triangleq$ (* initially, no resources have been requested or allocated *)

$$\wedge unsat = [Clients \rightarrow \{\}]$$

$$\wedge alloc = [Clients \rightarrow \{\}]$$

(* Client c requests set S of resources, provided it has no outstanding requests and no a

$Request(c,S) \triangleq$

$$\wedge unsat[c] = \{\} \wedge alloc[c] = \{\}$$

$$\wedge S \neq \{\} \wedge unsat' = [unsat \text{ EXCEPT } ![c] = S]$$

$$\wedge \text{UNCHANGED } alloc$$

(* Allocation of a set of available resources to a client that requested them. *)

$Allocate(c,S) \triangleq$

$$\wedge S \neq \{\} \wedge S \subseteq available \cap unsat[c]$$

$$\wedge alloc' = [alloc \text{ EXCEPT } ![c] = @ \cup S]$$

$$\wedge unsat' = [unsat \text{ EXCEPT } ![c] = @ \setminus S]$$

A first solution in TLA⁺ (3/4)

$$\begin{aligned} \text{Return}(c, S) &\triangleq \quad (* \text{ Client } c \text{ returns a set of resources that it holds. } *) \\ &\quad \wedge S \neq \{\} \wedge S \subseteq \text{alloc}[c] \\ &\quad \wedge \text{alloc}' = [\text{alloc} \text{ EXCEPT } ![c] = @ \setminus S] \\ &\quad \wedge \text{UNCHANGED } \text{unsat} \end{aligned}$$
$$\begin{aligned} \text{Next} &\triangleq \quad (* \text{ The next-state relation. } *) \\ &\quad \exists c \in \text{Clients}, S \in \text{SUBSET Resources} : \\ &\quad \quad \text{Request}(c, S) \vee \text{Allocate}(c, S) \vee \text{Return}(c, S) \end{aligned}$$
$$\text{vars} \triangleq \langle \text{unsat}, \text{alloc} \rangle$$

$$\begin{aligned} \text{SimpleAllocator} &\triangleq \quad (* \text{ The complete high-level specification. } *) \\ &\quad \wedge \text{Init} \wedge \Box [\text{Next}]_{\text{vars}} \\ &\quad \wedge \forall c \in \text{Clients} : \text{WF}_{\text{vars}}(\text{Return}(c, \text{alloc}[c])) \\ &\quad \wedge \forall c \in \text{Clients} : \text{SF}_{\text{vars}}(\exists S \in \text{SUBSET Resources} : \text{Allocate}(c, S)) \end{aligned}$$

A first solution in TLA⁺ (4/4)

(* Definition of system properties *)

ResourceMutex $\triangleq \forall c_1, c_2 \in \text{Clients} : c_1 \neq c_2 \Rightarrow \text{alloc}[c_1] \cap \text{alloc}[c_2] = \{\}$

ClientsWillObtain $\triangleq \forall c \in \text{Clients}, r \in \text{Resources} :$
 $(r \in \text{unsat}[c]) \leadsto (r \in \text{alloc}[c])$

ClientsWillReturn $\triangleq \forall c \in \text{Clients} : \text{unsat}[c] = \{\} \leadsto \text{alloc}[c] = \{\}$

InfOftenHappy $\triangleq \forall c \in \text{Clients} : \Box \Diamond (\text{unsat}[c] = \{\})$

THEOREM *SimpleAllocator* $\Rightarrow \Box \text{ResourceMutex}$

THEOREM *SimpleAllocator* $\Rightarrow \text{ClientsWillObtain}$

THEOREM *SimpleAllocator* $\Rightarrow \text{ClientsWillReturn}$

THEOREM *SimpleAllocator* $\Rightarrow \text{InfOftenHappy}$

We will later see how these theorems can be verified in TLA⁺

Summary

- specification language based on TLA and set theory
- TLA: stuttering invariant fragment of LTL
- specifications and properties represented as formulas
- TLA⁺: untyped ZF for specifying data structures
- module structure for hierarchical specifications
- system specification: explicit encoding of transition system
- TLA⁺ supports different specification styles

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺
- 4 Verification of TLA⁺ models**
 - Model checking finite instances
 - Deductive system verification
- 5 Structuring models: refinement and (de)composition

Verification and validation

- Validation: are we building the right system?
 - ▶ compare formal model against informal user requirements
 - ▶ **techniques:** animation, prototyping, testing, reviews
- Verification: are we building the system right?
 - ▶ use formal techniques to establish properties of a model
 - ▶ compare two formal models at different levels of abstraction
 - ▶ **techniques:** theorem proving, model and equivalence checking
- In the following: focus on verification

Verification by model checking and theorem proving

- Algorithmic verification – state space exploration $\mathcal{T} \models \Phi$
 - ▶ effectively construct state graph of transition system \mathcal{T}
 - ▶ use semantic definitions to evaluate property over \mathcal{T}
 - ▶ effective mechanization for certain properties over finite transition systems (and for some classes of infinite ones)
 - ▶ TLC: model checker for finite TLA⁺ models
- Deductive verification – theorem proving $\mathcal{T} \vdash \Phi$
 - ▶ define logic for expressing system properties
 - ▶ establish proof rules for relating system and properties
 - ▶ implement proof system within a theorem prover
 - ▶ TLAPS: interactive proof support for TLA⁺

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA^+
- 4 Verification of TLA^+ models**
 - **Model checking finite instances**
 - Deductive system verification
- 5 Structuring models: refinement and (de)composition

Definition

A state formula J is an **invariant** of a transition system $\mathcal{T} = (Q, I, \delta)$ if $\llbracket J \rrbracket_q = \text{tt}$ holds for every reachable state q .

- Invariants: fundamental correctness properties
 - ▶ exclusive access of processes to resources
 - ▶ non-corruption of messages
 - ▶ in-order delivery of messages
- Special case: inductive invariants
 - ▶ J is an **inductive invariant** if $\llbracket J \rrbracket_q = \text{tt}$ for every $q \in I$ and if for all $(q, q') \in \delta$, if $\llbracket J \rrbracket_q = \text{tt}$ then $\llbracket J \rrbracket_{q'} = \text{tt}$
 - ▶ **Note:** second condition required even for unreachable q
 - ▶ inductive invariants: basis of deductive system verification
 - ▶ some methods (such as B or Z) document inductive invariants

Invariant checking

- Idea of the algorithm

- ▶ enumerate reachable states of \mathcal{T} (finite state!)
- ▶ verify that all reachable states satisfy I

- Pseudo-code

```
boolean verifyInvariant(TransSystem ts, Predicate inv) {  
  Set visited = new Set();  
  Set todo = ts.getInitials();  
  while (!todo.isEmpty()) {  
    State s = todo.removeSomeElement()  
    if (!visited.contains(s)) {  
      if (!s.satisfies(inv))  
        return false;  
      visited.add(s);  
      todo.addAll(ts.getSuccessors(s));  
    }  
  }  
  return true;  
}
```

Invariant checking: implementation

- Depth-first search: `stack todo`
 - ▶ stack contains counter-example in case invariant does not hold
 - ▶ counter-examples may be unnecessarily long
- Breadth-first search: `queue todo`
 - ▶ ensures counter-examples of minimal length
 - ▶ counter-example generation needs additional back-pointers
- Problem: large state spaces (beyond 10^6 – 10^7 states)
 - ▶ use disk to store set `visited`
 - ▶ compress: store signatures instead of complete states
 \leadsto incomplete search in case of collisions
 - ▶ reduce number of states or transitions to explore

Beyond invariant checking

- Basic idea extends to more complicated properties
 - ▶ beyond mere reachability (of state violating invariant)
 - ▶ apply graph algorithms to determine validity of properties
- Example: “process A will eventually acquire resource”
 - ▶ enumerate strongly connected components of state graph (Tarjan)
 - ▶ check if in all states of SCC, process A tries to acquire resource without ever holding it
 - ▶ complexity linear in size of state graph
 - ▶ alternative: “double reachability” search (Courcoubetis et al. 1992)
- On-the-fly model checking
 - ▶ properties of runs represented by finite automata over ω -words: \mathcal{A} accepts state sequences that violate property
 - ▶ general graph algorithms for solving language emptiness of $\mathcal{T} \times \mathcal{A}$
 - ▶ many implementations, e.g. Spin: <http://spinroot.com>

Model checking using TLC

- Define finite instance for model checking

- ▶ instantiate system parameters to fixed values

`N = 5, Clients = {c1, c2, c3}, ...`

- ▶ indicate formula defining system specification

`SPECIFICATION SimpleAllocator`

- ▶ indicate properties to verify

`INVARIANTS TypeInvariant, ResourceMutex`

`TEMPORAL ClientsWillObtain, InfOftenHappy`

- ▶ easiest done by creating a “model” in the TLA⁺ toolbox

Model checking using TLC

- Define finite instance for model checking

- ▶ instantiate system parameters to fixed values

`N = 5, Clients = {c1, c2, c3}, ...`

- ▶ indicate formula defining system specification

`SPECIFICATION SimpleAllocator`

- ▶ indicate properties to verify

`INVARIANTS TypeInvariant, ResourceMutex`

`TEMPORAL ClientsWillObtain, InfOftenHappy`

- ▶ easiest done by creating a “model” in the TLA⁺ toolbox

- Some limitations of TLC (see Lamport’s book for details)

- ▶ all values must be finite (integers, finite sets, arrays, ...)
- ▶ specifications must be in standard form $Init \wedge \Box [Next]_v \wedge L$
- ▶ primed state variables must be “assigned to” at first occurrence

$x' = [x \text{ EXCEPT } ![3] = @ + 1]$

$v' \in 1..nProcs$

$x[3]' = x[3] + 1$

$1 \leq v' \wedge v' \leq nProcs$

SimpleAllocator.tla

SimpleAllocator

Model Overview Advanced Options Model Checking Results

Model Overview



What is the behavior spec?

☐ Initial predicate and next-state relation

Init:

Next:

☒ Temporal formula

SimpleAllocator

☐ No Behavior Spec

What to check?

☒ Deadlock

Invariants

Formulas true in every reachable state.

- ☒
- TypeInvariant
-
- ☒
- ResourceMutex

Add

Edit

Remove

Properties

Temporal formulas true for every possible behavior.

- ☒
- ClientsWillObtain
-
- ☒
- ClientsWillReturn
-
- ☒
- InfOftenHappy

Add

Edit

Remove

What is the model?

Specify the values of declared constants.

Clients <- [model value] {c1, c2, c3}

Resources <- [model value] {r1, r2, r3}

Edit

Advanced parts of the model: [Additional definitions,](#)
[State constraints,](#)[Definition override,](#)
[Additional model values.](#)

How to run?

37M of 81M



Spec Status : parsed

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA^+
- 4 Verification of TLA^+ models**
 - Model checking finite instances
 - Deductive system verification**
- 5 Structuring models: refinement and (de)composition

Principle of deductive verification

- System and properties represented as TLA formulas

system described by *Spec* has property *Prop*
iff
formula *Prop* holds of every run of *Spec*
iff
implication $Spec \Rightarrow Prop$ is valid

- System verification reduces to TLA provability
- No restriction to finite-state specifications
- Now: verification rules for standard correctness properties

Proving invariants

- Invariant: formula $\Box I$ for state predicate I
 - ▶ characterize the set of reachable states of a system
 - ▶ basis for proving further correctness properties
 - ▶ deductive verification relies on **inductive invariants**

- Basic proof rule: (INV1)
$$\frac{I \wedge Next \Rightarrow I' \quad I \wedge v' = v \Rightarrow I'}{I \wedge \Box[Next]_v \Rightarrow \Box I}$$

- ▶ every transition (stuttering or not) preserves I
 - ▶ thus, if I holds initially, it will hold forever

- Reduce temporal conclusion to non-temporal hypotheses

Invariant for the hour clock

```
MODULE HourClock
EXTENDS Naturals
VARIABLE hr

HCini     $\triangleq$   $hr \in 0..23$ 
HCnxt     $\triangleq$   $hr' = \text{IF } hr = 23 \text{ THEN } 0 \text{ ELSE } hr + 1$ 
HC        $\triangleq$   $HCini \wedge \Box[HCnxt]_{hr} \wedge WF_{hr}(HCnxt)$ 
```

- Prove $HC \Rightarrow \Box HCini$

- ▶ using (INV1) and the definition of HC we must prove

$$hr \in 0..23 \wedge HCnxt \Rightarrow hr' \in 0..23$$

$$hr \in 0..23 \wedge hr' = hr \Rightarrow hr' \in 0..23$$

- ▶ both implications are clearly valid

Finding inductive invariants

- Rule (INV1) can be used to prove inductive invariants

- ▶ in general, need to strengthen proposed invariant
- ▶ use derived proof rule

$$(INV) \quad \frac{Init \Rightarrow J \quad J \wedge (Next \vee v' = v) \Rightarrow J' \quad J \Rightarrow I}{Init \wedge \Box[Next]_v \Rightarrow \Box I}$$

- ▶ J : inductive invariant that implies I

- Inductive invariants

- ▶ finding inductive invariants is often difficult
- ▶ once found, proving them is easy (in principle)
- ▶ inductive invariants contain algorithmic idea: document them!
- ▶ there are systematic techniques for strengthening invariants

Example: strengthen invariant

```

————— MODULE SyncQueue —————
EXTENDS Sequences
VARIABLES i, q, o
|
SQInit  $\triangleq q = \langle \rangle \wedge i = o$ 
SQEnq  $\triangleq i' \neq i \wedge q' = \text{Append}(q, i') \wedge o' = o$ 
SQDeq  $\triangleq q \neq \langle \rangle \wedge o' = \text{Head}(q) \wedge q' = \text{Tail}(q) \wedge i' = i$ 
SQLive  $\triangleq \text{WF}_{i, q, o}(\text{SQDeq})$ 
SQSpec  $\triangleq \text{SQInit} \wedge \Box[\text{SQEnq} \vee \text{SQDeq}]_{i, q, o} \wedge \text{SQLive}$ 

```

- Elements are enqueued only if input changes
- Therefore any two consecutive queue elements are different
- **Exercise:** formally prove

$$\text{SQSpec} \Rightarrow \Box(\forall i \in (1..Len(q) - 1) : q[i] \neq q[i + 1])$$

Excursion: weakest preconditions (1)

- For an action A and a state predicate P define

$$\mathbf{wp}(P, A) \triangleq \forall v'_1, \dots, v'_n : A \Rightarrow P' \quad (\equiv \neg \text{ENABLED}(A \wedge \neg P'))$$

where v'_1, \dots, v'_n are all free primed variables in A or P'

- $\mathbf{wp}(P, A)$: weakest precondition of P w.r.t. A

► true in those states all of whose A -successors satisfy P

- Examples

$$\begin{aligned} \mathbf{wp}(x = 5, x' = x + 1) &\equiv \forall x' : x' = x + 1 \Rightarrow x' = 5 \\ &\equiv x = 4 \end{aligned}$$

$$\begin{aligned} \mathbf{wp}(y \in S, S' = S \cup T \wedge y' = y) &\equiv \forall y', S' : S' = S \cup T \wedge y' = y \Rightarrow y' \in S' \\ &\equiv y \in S \cup T \end{aligned}$$

$$\begin{aligned} \mathbf{wp}(x > 0, x' = 0) &\equiv \forall x' : x' = 0 \Rightarrow x' > 0 \\ &\equiv \text{ff} \end{aligned}$$

Excursion: weakest preconditions (2)

- Rule (INV) can be rewritten as follows

$$\frac{Init \Rightarrow J \quad J \Rightarrow \mathbf{wp}(J, Next \vee v' = v) \quad J \Rightarrow I}{Init \wedge \Box[Next]_v \Rightarrow \Box I}$$

- Following heuristic can help finding inductive invariants

- 1 start with invariant to prove: $J := I$
- 2 try proving $J \wedge A \Rightarrow J'$ for each disjunct A of $[Next]_v$
if this proof fails, set $J := J \wedge \mathbf{wp}(J, A)$
- 3 repeat step 2 until
 - ★ either all proofs succeed and $Init \Rightarrow J$ holds:
 $\Rightarrow J$ is an inductive invariant
 - ★ or J is not implied by $Init$:
 $\Rightarrow I$ is not an invariant

- Heuristic need not terminate: generalize candidate invariant

Liveness 1: use fairness conditions

- Fairness conditions ensure that actions occur eventually
- Following rule derives liveness from weak fairness

$$\text{(WF1)} \quad \frac{\begin{array}{l} P \wedge [Next]_v \Rightarrow P' \vee Q' \\ P \wedge \langle Next \wedge A \rangle_v \Rightarrow Q' \\ P \Rightarrow \text{ENABLED } \langle A \rangle_v \end{array}}{\Box[Next]_v \wedge \text{WF}_v(A) \Rightarrow (P \leadsto Q)}$$

- hypotheses of (WF1) are again non-temporal

Examples

- Fairness of $HCnxt$ ensures that clock keeps ticking

$$HC \Rightarrow \forall k \in 0..23 : hr = k \leadsto hr = (k + 1)\%24$$

- ▶ by (WF1) and predicate logic it suffices to show

$$k \in 0..23 \wedge hr = k \wedge [HCnxt]_{hr} \Rightarrow hr' = k \vee hr' = (k + 1)\%24$$

$$k \in 0..23 \wedge hr = k \wedge \langle HCnxt \rangle_{hr} \Rightarrow hr' = (k + 1)\%24$$

$$k \in 0..23 \wedge hr = k \Rightarrow \text{ENABLED } \langle HCnxt \rangle_{hr}$$

- ▶ these formulas are again valid

- **Exercise:** show that elements advance in queue

$$SQSpec \Rightarrow \forall k \in 1..Len(q) : \forall x : q[k] = x \leadsto (o = x \vee q[k - 1] = x)$$

Correctness of rule (WF1)

Assumptions

- 1 $\sigma \models \Box[Next]_v \wedge WF_v(A)$ where $\sigma = s_0 s_1 \dots$
- 2 To prove that $\sigma \models P \leadsto Q$ assume that $\llbracket P \rrbracket_{s_n} = \mathbf{tt}$ for some $n \in \mathbb{N}$
- 3 Assume also that $\llbracket Q \rrbracket_{s_m} = \mathbf{ff}$ for all $m \geq n$

Show: contradiction

- 1 $\llbracket P \rrbracket_{s_m} = \mathbf{tt}$ for all $m \geq n$
[induction on $m \geq n$ using A1, A2, and hypothesis $P \wedge [Next]_v \Rightarrow P' \vee Q'$]
- 2 $\llbracket \text{ENABLED } \langle A \rangle_v \rrbracket_{s_m} = \mathbf{tt}$ for all $m \geq n$
[from (1) and hypothesis $P \Rightarrow \text{ENABLED } \langle A \rangle_v$]
- 3 $\llbracket \langle A \rangle_v \rrbracket_{s_m, s_{m+1}} = \mathbf{tt}$ for some $m \geq n$
[from (2) and A1: $WF_v(A)$]
- 4 $\llbracket Q \rrbracket_{s_{m+1}} = \mathbf{tt}$ for some $m \geq n$
[from (3), (1), A1, and hypothesis $P \wedge \langle Next \wedge A \rangle_v \Rightarrow Q'$]
- 5 Q.E.D. (contradiction with A3)

Liveness from strong fairness

- For (WF1), the action $\langle A \rangle_v$ must remain enabled forever
 - ▶ strong fairness only requires “infinitely often enabled”
 - ▶ appropriate proof rule

$$\begin{array}{c} P \wedge [Next]_v \Rightarrow P' \vee Q' \\ P \wedge \langle Next \wedge A \rangle_v \Rightarrow Q' \\ \text{(SF1)} \quad \frac{\Box P \wedge \Box [Next]_v \wedge \Box F \Rightarrow \Diamond \text{ENABLED } \langle A \rangle_v}{\Box [Next]_v \wedge \text{SF}_v(A) \wedge \Box F \Rightarrow (P \leadsto Q)} \end{array}$$

- Observations

- ▶ first two hypotheses are as in (WF1)
- ▶ third hypothesis is a temporal formula: F can be a conjunction of
 - fairness conditions: $\text{WF}_v(B) \equiv \Box \text{WF}_v(B)$, $\text{SF}_v(B) \equiv \Box \text{SF}_v(B)$
 - auxiliary “leadsto” formulas, invariants, ...

Complex liveness properties

- Rules (WF1) and (SF1) prove elementary liveness properties

- ▶ clock eventually displays next hour
- ▶ elements in queue will advance by one position

- How can we prove more complex properties?

- ▶ clock will eventually display 12 o'clock

$$HC \Rightarrow \Box \Diamond (hr = 12)$$

- ▶ enqueued values will eventually be output

$$SQSpec \Rightarrow ((\exists k \in 1..Len(q) : q[k] = x) \leadsto o = x)$$

- Informal argument: repeat elementary liveness property

- ▶ every tick of the clock brings us closer to noon
- ▶ every output action moves the element closer to the output channel

Definition

A binary relation $\prec \subseteq D \times D$ is **well-founded** if there is no infinite descending chain $d_0 \succ d_1 \succ d_2 \succ \dots$ of elements $d_i \in D$.

• Observations

- ▶ well-founded relations are irreflexive and asymmetric
- ▶ every non-empty subset of D contains a minimal element

• Examples

- ▶ $<$ is well-founded over \mathbb{N} (also over ordinal numbers)
- ▶ lexicographic ordering on sequences over bounded size
- ▶ lexicographic ordering on $\text{Seq}(\{a, b\})$ is not well-founded:

$$b \succ ab \succ aab \succ aaab \succ \dots$$

Liveness from well-founded relations

- Additional proof rule

$$\text{(WFO)} \quad \frac{(D, \prec) \text{ well-founded} \quad F \wedge d \in D \Rightarrow (H(d) \leadsto G \vee (\exists e \in D : e \prec d \wedge H(e)))}{F \Rightarrow ((\exists d \in D : H(d)) \leadsto G)} \quad (d \text{ does not occur in } G)$$

- Observations

- ▶ must prove another “leadsto” property, typically based on fairness
- ▶ first premise should be verified by reasoning about data

- Exercise: prove the correctness of rule (WFO)

Example

Prove $HC \Rightarrow ((\exists d \in 0..23 : hr = d) \leadsto hr = 12)$

- Define the well-founded relation \prec on $0..23$ by

$$\begin{aligned} dist(d) &\triangleq \text{IF } d \leq 12 \text{ THEN } 12 - d \text{ ELSE } 36 - d \quad (* \text{ distance from } 12 *) \\ d \prec e &\triangleq dist(d) < dist(e) \end{aligned}$$

- Applying (WFO) we have to prove

$$HC \wedge d \in 0..23 \Rightarrow (hr = d \leadsto (hr = 12 \vee \exists e \in 0..23 : e \prec d \wedge hr = e))$$

- This follows from the formula

$$HC \Rightarrow \forall k \in (0..23) : hr = k \leadsto hr = (k + 1) \bmod 24$$

shown earlier, using the definition of \prec

Combining properties: simple temporal logic

Application of verification rules is supported by the laws of

- first-order logic,
- theories formalizing the data (in particular, set theory),
- and laws of temporal logic such as

$$\text{(STL1)} \quad \frac{F}{\Box F}$$

$$\text{(STL2)} \quad \Box F \Rightarrow F$$

$$\text{(STL3)} \quad \Box \Box F \equiv \Box F$$

$$\text{(STL4)} \quad \Box(F \Rightarrow G) \Rightarrow (\Box F \Rightarrow \Box G)$$

$$\text{(STL5)} \quad \Box(F \wedge G) \equiv (\Box F \wedge \Box G)$$

$$\text{(STL6)} \quad \Diamond \Box(F \wedge G) \equiv \Diamond \Box F \wedge \Diamond \Box G$$

$$\text{(TLA1)} \quad \frac{P \wedge t' = t \Rightarrow P'}{\Box P \equiv P \wedge \Box[P \Rightarrow P']_t}$$

$$\text{(TLA2)} \quad \frac{I \wedge I' \wedge [A]_t \Rightarrow [B]_u}{\Box I \wedge \Box[A]_t \Rightarrow \Box[B]_u}$$

Validity of propositional temporal logic is mechanically decidable

Proof support: the TLA⁺ proof system

- Assist user in writing proofs

- ▶ hierarchical, declarative proof language
- ▶ skeleton of invariant proof

THEOREM $Spec \Rightarrow \Box Inv$

⟨1⟩1. $Init \Rightarrow Inv$

⟨1⟩2. ASSUME $Inv, Next$

PROVE Inv'

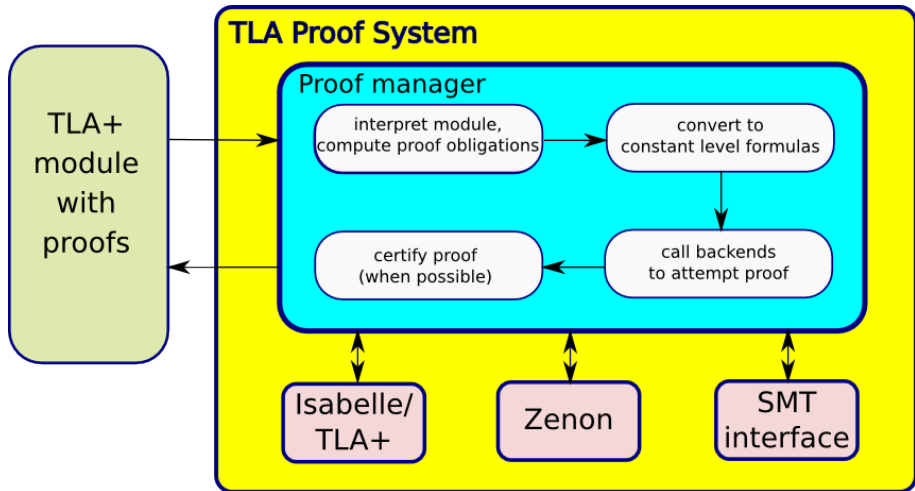
⟨1⟩3. QED

BY ⟨1⟩1, ⟨1⟩2, INV1 DEF $Spec$

- ▶ proofs of steps ⟨1⟩1 and ⟨1⟩2 to be filled in

- Verification of leaf proofs by theorem provers
- Philosophy: decompose proof steps until they become “trivial”
- First release: <http://msr-inria.inria.fr/~doligez/tlaps/>

Architecture of the TLAPS



Summary

- model checking: push-button verification of finite instances
- deductive approach: reduce system verification to logical validity
- proof rules for invariants and for liveness properties
- TLA: quickly reduce to non-temporal verification conditions
- machine assistance: theorem proving through the TLAPS

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition**
 - Refinement of specifications
 - Hiding (encapsulation) of internal state
 - Composition and decomposition

Modeling “in the large”

- So far: specifications at single level of abstraction
- This chapter:
 - ▶ compare models at different levels of abstraction: refinement
 - ▶ hiding (encapsulation) of internal state
 - ▶ composition from and decomposition into system component
- These concepts are represented in TLA⁺ by logical connectives

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition**
 - **Refinement of specifications**
 - Hiding (encapsulation) of internal state
 - Composition and decomposition

Reminder: refining the hour clock

```

┌────────────────── MODULE HourMinuteClock ───────────────────┐
| EXTENDS Naturals, HourClock                                   |
| VARIABLE min                                                  |
|──────────────────────────────────────────────────────────────────|
| HMCini       $\triangleq$  HCini  $\wedge$  min  $\in$  0..59                |
| Min          $\triangleq$  min' = IF min = 59 THEN 0 ELSE min + 1    |
| Hr           $\triangleq$  IF min = 59 THEN HCnxt ELSE hr' = hr      |
| HMCnxt       $\triangleq$  Min  $\wedge$  Hr                                |
| HMC          $\triangleq$  HMCini  $\wedge$   $\Box$ [HMCnxt] $\langle hr, min \rangle$   $\wedge$  WF $\langle hr, min \rangle$ (HMCnxt) |
|──────────────────────────────────────────────────────────────────|
| THEOREM HMC  $\Rightarrow$  HC                                         |
└──────────────────────────────────────────────────────────────────┘

```

- Every run of *HMC* also satisfies the hour-clock specification
- Stuttering invariance essential for allowing minute ticks

Proving the implication $HMC \Rightarrow HC$


- Initial condition $HMCini \Rightarrow HCini$
 - ▶ obvious from definition of $HMCini$
- Next-state relation $HMC \Rightarrow \Box[HCnxt]_{hr}$
 - ▶ easy by definition of $HMCnxt$, formally supported by rule (TLA2)
- Fairness condition $HMC \Rightarrow WF_{hr}(HCnxt)$
 - ▶ unfold definition of $WF_{hr}(HCnxt)$ or use derived proof rule

$$(WF2) \quad \frac{\begin{array}{c} \langle N \wedge P \wedge A \rangle_v \Rightarrow \langle B \rangle_w \\ P \wedge \text{ENABLED } \langle B \rangle_w \Rightarrow \text{ENABLED } \langle A \rangle_v \\ \Box[N \wedge [\neg B]_w]_v \wedge WF_v(A) \wedge \Box F \wedge \Diamond \Box \text{ENABLED } \langle B \rangle_w \Rightarrow \Diamond \Box P \end{array}}{\Box[N]_v \wedge WF_v(A) \wedge \Box F \Rightarrow WF_w(B)}$$

- **Exercise:** verify the refinement using TLC


Refinement as implication (trace inclusion)

- Refinement may add state variables (“implementation detail”)
- High-level actions implemented by several low-level steps
 - ▶ actions except last one do not affect high-level state
 - ▶ final action corresponds to high-level effect



Refinement as implication (trace inclusion)

- Refinement may add state variables (“implementation detail”)
- High-level actions implemented by several low-level steps
 - ▶ actions except last one do not affect high-level state
 - ▶ final action corresponds to high-level effect



- Stuttering invariance is crucial for refinement as implication
 - ▶ must also preserve high-level fairness conditions
 - ▶ implementation inherits all high-level properties
- Implementation may reduce non-determinism
 - ▶ branching structure of high-level model need not be preserved

Resource allocator revisited

- Consider the following scenario
 - ▶ clients c_1 and c_2 request resources r_1 and r_2
 - ▶ allocator gives r_1 to c_1 and r_2 to c_2
- Is the system deadlocked?
- Why didn't we see the problem when running TLC?

Resource allocator revisited

- Consider the following scenario
 - ▶ clients c_1 and c_2 request resources r_1 and r_2
 - ▶ allocator gives r_1 to c_1 and r_2 to c_2
- Is the system deadlocked?
 - ▶ **No:** c_1 may return the resource r_1 , unblocking the system
- Why didn't we see the problem when running TLC?
 - ▶ in fact, c_1 must eventually return r_1 due to the fairness condition
$$\forall c \in Clients : WF_{vars}(\text{Return}(c, \text{alloc}[c]))$$
 - ▶ therefore, all required liveness properties hold in the model

Resource allocator revisited

- Consider the following scenario
 - ▶ clients c_1 and c_2 request resources r_1 and r_2
 - ▶ allocator gives r_1 to c_1 and r_2 to c_2
- Is the system deadlocked?
 - ▶ **No:** c_1 may return the resource r_1 , unblocking the system
- Why didn't we see the problem when running TLC?
 - ▶ in fact, c_1 must eventually return r_1 due to the fairness condition
$$\forall c \in Clients : WF_{vars}(\text{Return}(c, \text{alloc}[c]))$$
 - ▶ therefore, all required liveness properties hold in the model
- **The specification *SimpleAllocator* is wrong!**

Weakening the fairness condition

- Intended fairness requirement (cf. informal description)

- ▶ clients must eventually return all resources it holds
after having received all resources it requested
- ▶ this can be expressed in TLA⁺ as

$$\forall c \in Clients : WF_{vars}(\text{unsat}[c] = \{\} \wedge \text{Return}(c, \text{alloc}[c]))$$

- With this weaker condition the liveness properties no longer hold

- ▶ TLC reports the expected counter-example

Weakening the fairness condition

- Intended fairness requirement (cf. informal description)

- ▶ clients must eventually return all resources it holds
after having received all resources it requested
- ▶ this can be expressed in TLA⁺ as

$$\forall c \in Clients : WF_{vars}(unsat[c] = \{\} \wedge Return(c, alloc[c]))$$

- With this weaker condition the liveness properties no longer hold

- ▶ TLC reports the expected counter-example

- Lessons learnt

- 1 specifying fairness can be tricky
- 2 verification is not a panacea: specifications need to be validated

- How can we correct the specification?

Resource allocator: second solution

- Idea: allocator maintains a schedule of clients to satisfy
 - ▶ ensure that all requests can be satisfied, even in the worst case
- Principle of operation
 - ▶ allocate resource r to client c only if c appears in the schedule and no client c' that is scheduled before c requests r
 - ▶ upon receiving a new request from client c , put c into a pool of clients awaiting scheduling
 - ▶ eventually append clients from the pool to the schedule (in arbitrary order: can be refined later)
- In the following: formulation in TLA⁺

Scheduling allocator in TLA⁺ (1/5)

MODULE *SchedulingAllocator*

EXTENDS *FiniteSets*, *Sequences*, *Naturals*

CONSTANTS *Clients*, *Resources*

ASSUME *IsFiniteSet*(*Resources*)

VARIABLES

unsat, (* *unsat*[*c*] denotes the outstanding requests of client *c* *)

alloc, (* *alloc*[*c*] denotes the resources allocated to client *c* *)

pool, (* clients with unsatisfied requests that have not been scheduled *)

sched (* schedule (sequence of clients) *)

TypeInvariant \triangleq

$\wedge \textit{unsat} \in [\textit{Clients} \rightarrow \text{SUBSET } \textit{Resources}]$

$\wedge \textit{alloc} \in [\textit{Clients} \rightarrow \text{SUBSET } \textit{Resources}]$

$\wedge \textit{pool} \in \text{SUBSET } \textit{Clients}$

$\wedge \textit{sched} \in \text{Seq}(\textit{Clients})$

Scheduling allocator in TLA⁺ (2/5)

$PermSeqs(S) \triangleq$ (* Permutation sequences of finite set S. *)

LET $perms[ss \in SUBSET S] \triangleq$

 IF $ss = \{\}$ THEN $\langle \rangle$

 ELSE LET $ps \triangleq [x \in ss \mapsto \{Append(sq, x) : sq \in perms[ss \setminus \{x\}]]$

 IN UNION $\{ps[x] : x \in ss\}$

IN $perms[S]$

$Drop(seq, i) \triangleq$ (* remove element at position i from sequence seq *)

$SubSeq(seq, 1, i - 1) \circ SubSeq(seq, i + 1, Len(seq))$

$available \triangleq$ (* set of resources free for allocation *)

$Resources \setminus (UNION \{alloc[c] : c \in Clients\})$

Scheduling allocator in TLA⁺ (3/5)

$Init \triangleq$

$\wedge \text{unsat} = [\text{Clients} \rightarrow \{\}] \wedge \text{alloc} = [\text{Clients} \rightarrow \{\}]$

$\wedge \text{pool} = \{\} \wedge \text{sched} = \langle \rangle$

$\text{Request}(c, S) \triangleq$ (* Client c requests set S of resources. *)

$\wedge \text{unsat}[c] = \{\} \wedge \text{alloc}[c] = \{\}$

$\wedge S \neq \{\} \wedge \text{unsat}' = [\text{unsat} \text{ EXCEPT } ![c] = S]$

$\wedge \text{pool}' = \text{pool} \cup \{c\}$

$\wedge \text{UNCHANGED } \langle \text{alloc}, \text{sched} \rangle$

$\text{Return}(c, S) \triangleq$ (* Client c returns a set of resources that it holds. *)

$\wedge S \neq \{\} \wedge S \subseteq \text{alloc}[c]$

$\wedge \text{alloc}' = [\text{alloc} \text{ EXCEPT } ![c] = @ \setminus S]$

$\wedge \text{UNCHANGED } \langle \text{unsat}, \text{pool}, \text{sched} \rangle$

Scheduling allocator in TLA⁺ (4/5)

$Schedule \triangleq$ (* Extend allocator schedule by the processes from the pool. *)
 $\wedge pool \neq \{\}$
 $\wedge \exists sq \in PermSeqs(pool) : sched' = sched \circ sq$
 $\wedge pool' = \{\}$
 $\wedge UNCHANGED \langle unsat, alloc \rangle$

$Allocate(c, S) \triangleq$ (* Allocate set S of available resources to client c . *)
 $\wedge S \neq \{\} \wedge S \subseteq available \cap unsat[c]$
 $\wedge \exists i \in 1..Len(sched) :$
 $\wedge sched[i] = c$
 $\wedge \forall j \in 1..i-1 : unsat[sched[j]] \cap S = \{\}$
 $\wedge sched' = \text{IF } S = unsat[c] \text{ THEN Drop}(sched, i) \text{ ELSE } sched$
 $\wedge alloc' = [alloc \text{ EXCEPT } ![c] = @ \cup S]$
 $\wedge unsat' = [unsat \text{ EXCEPT } ![c] = @ \setminus S]$
 $\wedge UNCHANGED pool$

Scheduling allocator in TLA⁺ (5/5)

$Next \triangleq$ (* The next-state relation. *)
 $\vee \exists c \in Clients, S \in \text{SUBSET Resources} :$
 $Request(c, S) \vee Allocate(c, S) \vee Return(c, S)$
 $\vee \textit{Schedule}$

$vars \triangleq \langle unsat, alloc, pool, sched \rangle$

$SchedulingAllocator \triangleq$
 $\wedge Init \wedge \Box [Next]_{vars}$
 $\wedge \forall c \in Clients : WF_{vars}(unsat[c] = \{\} \wedge Return(c, alloc[c]))$
 $\wedge \forall c \in Clients : \textcolor{blue}{WF}_{vars}(\exists S \in \text{SUBSET Resources} : Allocate(c, S))$
 $\wedge \textcolor{blue}{WF}_{vars}(\textit{Schedule})$

Verifying the scheduling allocator

- Scheduling allocator satisfies all correctness requirements
- Crucial invariant
 - ▶ full request of any scheduled client can be satisfied from the resources that will be available after all previously scheduled clients released all resources they held or requested

$$\begin{aligned} \forall i \in 1..Len(sched) : unsat[sched[i]] \subseteq \\ \quad available \\ \cup \text{UNION } \{unsat[sched[j]] \cup alloc[sched[j]] : j \in 1..i-1\} \\ \cup \text{UNION } \{alloc[c] : c \in Clients\} \end{aligned}$$

Verifying the scheduling allocator

- Scheduling allocator satisfies all correctness requirements
- Crucial invariant
 - ▶ full request of any scheduled client can be satisfied from the resources that will be available after all previously scheduled clients released all resources they held or requested

$$\begin{aligned} \forall i \in 1..Len(sched) : unsat[sched[i]] \subseteq \\ \quad available \\ \cup \text{UNION } \{unsat[sched[j]] \cup alloc[sched[j]] : j \in 1..i-1\} \\ \cup \text{UNION } \{alloc[c] : c \in Clients\} \end{aligned}$$

- Successful verification using TLC (for small instances)
- Do you believe that the model is now correct?

Scheduling allocator as a refinement

- In fact, *SchedulingAllocator* refines *SimpleAllocator*

```
MODULE AllocatorRefinement  
EXTENDS SchedulingAllocator  
Simple  $\triangleq$  INSTANCE SimpleAllocator  
THEOREM SchedulingAllocator  $\Rightarrow$  Simple!SimpleAllocator
```

- ▶ the weaker fairness conditions of the scheduling allocator suffice thanks to scheduling
 - ▶ the allocator does no longer require the cooperation of the clients for ensuring that they will eventually receive all resources
- Correctness properties are inherited by *SchedulingAllocator*

Outline

- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition**
 - Refinement of specifications
 - Hiding (encapsulation) of internal state**
 - Composition and decomposition

Reminder: specification of a FIFO component

MODULE *SyncQueue*

EXTENDS *Sequences*

VARIABLES i, q, o


$SQInit \triangleq q = \langle \rangle \wedge i = o$

$SQEnq \triangleq i' \neq i \wedge q' = Append(q, i') \wedge o' = o$

$SQDeq \triangleq q \neq \langle \rangle \wedge o' = Head(q) \wedge q' = Tail(q) \wedge i' = i$

$SQLive \triangleq WF_{i,q,o}(SQDeq)$

$SQSpec \triangleq SQInit \wedge \Box [SQEnq \vee SQDeq]_{i,q,o} \wedge SQLive$



- The internal queue q is an “implementation detail”

- ▶ it should not be part of the interface
- ▶ FIFO component should function **as if there were** a queue
- ▶ do not expose implementation detail, especially for refinement

Hiding in TLA⁺

- TLA uses existential quantification to represent hiding

```

      _____ MODULE FIFO _____
      |
      | VARIABLES i,o
      | Queue(q)  $\triangleq$  INSTANCE SyncQueue
      |
      | _____
      | FIFO  $\triangleq$   $\exists q : \textit{Queue}(q)!\textit{SQSpec}$ 
      |

```

- ▶ module *FIFO* contains the “external” specification of the queue
- ▶ separate interface and implementation

Hiding in TLA⁺

- TLA uses existential quantification to represent hiding

```

┌────────── MODULE FIFO ─────────┐
│  VARIABLES i,o                    │
│  Queue(q)  $\triangleq$  INSTANCE SyncQueue │
├──────────────────────────────────┤
│  FIFO  $\triangleq$   $\exists q : \text{Queue}(q)!SQSpec$  │
└──────────────────────────────────┘

```

- ▶ module *FIFO* contains the “external” specification of the queue
 - ▶ separate interface and implementation
- Extend syntax of TLA formulas
 - ▶ if F is a formula and v a state variable then $\exists v : F$ is a formula
 - ▶ **intuitive meaning:** $\sigma \models \exists v : F$ if $\tau \models F$ for some τ that differs from σ only in the valuations of v

Naive semantics of quantification

- First attempt at defining semantics

$$\sigma, \xi \models \exists v : F \quad \text{iff} \quad \tau, \xi \models F \text{ for some } \tau \text{ s.t. } \tau_n =_v \sigma_n \text{ for all } n \in \mathbb{N}$$

Naive semantics of quantification

- First attempt at defining semantics

$$\sigma, \xi \models \exists v : F \quad \text{iff} \quad \tau, \xi \models F \quad \text{for some } \tau \text{ s.t. } \tau_n =_v \sigma_n \text{ for all } n \in \mathbb{N}$$

- Consider $F \triangleq v = 0 \wedge \Box[v = 1]_w$

τ				→
v	0	1	1	...
w	0	0	1	...

σ				→
v	0	0	1	...
w	0	0	1	...

- ▶ F essentially asserts that v has changed value before w changes
- ▶ clearly, $\tau \models F$, and therefore $\sigma \models \exists v : F$
- ▶ however, $\exists v : F$ would not hold of σ with second state removed

- Violation of stuttering invariance

Formal semantics of quantification

- Solution: define semantics with stuttering invariance “built in”

$\sigma, \xi \models \exists v : F$ iff there exist $\rho \approx \sigma$ and $\tau =_v \rho$ s.t. $\tau, \xi \models F$

- ▶ both σ and its variant satisfy $\exists v : v = 0 \wedge \Box[v = 1]_w$
- ▶ in fact, this formula is valid

Formal semantics of quantification

- Solution: define semantics with stuttering invariance “built in”

$\sigma, \xi \models \exists v : F$ iff there exist $\rho \approx \sigma$ and $\tau =_v \rho$ s.t. $\tau, \xi \models F$

- ▶ both σ and its variant satisfy $\exists v : v = 0 \wedge \Box[v = 1]_w$
- ▶ in fact, this formula is valid

- Clock example: have $HC \Rightarrow \exists min : HMC$

- ▶ implementation of hour clock may contain hidden minute hand

- Verification: standard proof rules are sound

(\exists -I) $F(t) \Rightarrow \exists v : F(v)$ (t state function: “refinement mapping”)


(\exists -E) $\frac{F \Rightarrow G}{(\exists v : F) \Rightarrow G}$ (v not free in G)

- ▶ for completeness, need more introduction rules
- ▶ TLC does not support \exists : provide explicit refinement mapping

Outline


- 1 Motivation and introductory example
- 2 Transition systems and their properties
- 3 System Specification in TLA⁺
- 4 Verification of TLA⁺ models
- 5 Structuring models: refinement and (de)composition**
 - Refinement of specifications
 - Hiding (encapsulation) of internal state
 - Composition and decomposition**

Parallel composition of components




- Assume: components A and B modeled by $ASpec$ and $BSpec$
 - ▶ interface represented by the shared variables $a2b$ and $b2a$
 - ▶ both specifications accommodate changes due to other component
 - ▶ assume: internal variables hidden (or renamed to be disjoint)

Parallel composition of components



- Assume: components A and B modeled by $ASpec$ and $BSpec$
 - ▶ interface represented by the shared variables $a2b$ and $b2a$
 - ▶ both specifications accommodate changes due to other component
 - ▶ assume: internal variables hidden (or renamed to be disjoint)
- System containing both components must satisfy $ASpec \wedge BSpec$
 - ▶ stuttering invariance is again important here
- System model may require additional constraints
 - ▶ typical example: interleaving assumption

Example: composition of two FIFO channels



MODULE *TwoQueues*

VARIABLES i, m, o

$LeftFIFO \triangleq$ INSTANCE *FIFO* WITH $o \leftarrow m$


$RightFIFO \triangleq$ INSTANCE *FIFO* WITH $i \leftarrow m$

$LongFIFO \triangleq LeftFIFO!FIFO \wedge RightFIFO!FIFO$

$Single \triangleq$ INSTANCE *FIFO*

THEOREM $LongFIFO \Rightarrow Single!FIFO$??

Example: composition of two FIFO channels



```

MODULE TwoQueues
  VARIABLES  $i, m, o$ 
  LeftFIFO  $\triangleq$  INSTANCE FIFO WITH  $o \leftarrow m$ 
  RightFIFO  $\triangleq$  INSTANCE FIFO WITH  $i \leftarrow m$ 
  LongFIFO  $\triangleq$  LeftFIFO!FIFO  $\wedge$  RightFIFO!FIFO
  Single  $\triangleq$  INSTANCE FIFO
  THEOREM LongFIFO  $\Rightarrow$  Single!FIFO ??

```

- *LongFIFO* allows for simultaneous enqueueing and dequeueing
 - ▶ *Single!FIFO* does not, due to interleaving assumption
 - ▶ assert interleaving when composing

Interleaving composition

- Interleaving composition: add a “synchronizer”

THEOREM $LongFIFO \wedge \Box[i' = i \vee o' = o]_{i,o} \Rightarrow Spec!FIFO$ where

$$\begin{aligned} LongFIFO &\equiv \wedge \exists q : SyncQueue(i, q, m)!SQSpec \\ &\quad \wedge \exists q : SyncQueue(m, q, o)!SQSpec \\ Spec!FIFO &\equiv \exists q : SyncQueue(i, q, o)!SQSpec \end{aligned}$$

Interleaving composition

- Interleaving composition: add a “synchronizer”

THEOREM $LongFIFO \wedge \Box[i' = i \vee o' = o]_{i,o} \Rightarrow Spec!FIFO$ where

$$\begin{aligned} LongFIFO &\equiv \wedge \exists q : SyncQueue(i, q, m)!SQSpec \\ &\quad \wedge \exists q : SyncQueue(m, q, o)!SQSpec \\ Spec!FIFO &\equiv \exists q : SyncQueue(i, q, o)!SQSpec \end{aligned}$$

- Proof (outline)

- ▶ using rules $(\exists\text{-E})$ and $(\exists\text{-I})$ it is enough to prove

$$\begin{aligned} &\wedge SyncQueue(i, q_2, m)!SQSpec \wedge SyncQueue(m, q_1, o)!SQSpec \\ &\wedge \Box[i' = i \vee o' = o]_{i,o} \\ &\Rightarrow SyncQueue(i, t, o)!SQSpec \end{aligned}$$

for some state function t . Choose $t \triangleq q_1 \circ q_2$.

- Exercise: formally carry out this proof

Decomposition of system specification

- Obtain component specifications from high-level system model
 - ▶ more in line with refinement approach to system development
 - ▶ technically often simpler than composing independent components
- Partition system state & identify component responsibilities
 - ▶ assume that no variable can be written by two different components
 - ▶ variables can be read by any number of components
 - ▶ generalization possible, but technically more involved
- Example: development of a producer-consumer system

High-level specification



- First abstraction: only specify the buffer
 - ▶ ignore nature of products and how they are produced/consumed
 - ▶ do not specify communication between buffer and environment
- Assume FIFO buffer of fixed capacity

High-level specification in TLA⁺

MODULE *AbstractPC*

EXTENDS *Sequences*

CONSTANTS *Product, None, buffCap*

VARIABLES *buffer, in, out*

Init \triangleq *buffer* = $\langle \rangle$ \wedge *in* = *None* \wedge *out* = *None*

Get \triangleq \wedge *Len(buffer)* < *buffCap*
 \wedge *buffer'* = *Append(buffer, in)*
 \wedge UNCHANGED \langle *in, out* \rangle

Put \triangleq \wedge *Len(buffer)* > 0
 \wedge *out'* = *Head(buffer)* \wedge *buffer'* = *Tail(buffer)*
 \wedge UNCHANGED *in*

vars \triangleq \langle *in, buffer, out* \rangle

Spec \triangleq *Init* \wedge \Box [*Get* \vee *Put*]_{*vars*} \wedge WF_{*vars*}(*Put*)

THEOREM *Spec* \Rightarrow $\Box(\Diamond\langle$ *Get* $\rangle_{vars} \Rightarrow \Diamond\langle$ *Put* $\rangle_{vars})$


Refinement: introducing communication



- Explicitly model communication between components
 - ▶ assume asynchronous communication for distributed system
- Synchronize sender and receiver over channel
 - ▶ signal availability of new value and successful reception
 - ▶ in particular, block producer when the buffer is full
 - ▶ avoid variables written by two different components

Channel specification in TLA⁺

- Generic channel specification with “split bit” handshake



- ▶ channel content represented by ch , control bits tmt and rcv
- ▶ channel empty (i.e., value may be sent) when $tmt = rcv$
- ▶ channel full (i.e., value may be received) when $tmt \neq rcv$

```

MODULE Channel
CONSTANTS Value, NoValue
VARIABLES ch, tmt, rcv
Init      ≜ ch = NoValue ∧ tmt = FALSE ∧ rcv = FALSE
Snd(v)    ≜ tmt = rcv ∧ ch' = v ∧ tmt' = ¬tmt ∧ rcv' = rcv
Rcv       ≜ tmt ≠ rcv ∧ rcv' = ¬rcv ∧ UNCHANGED ⟨ch, tmt⟩
Next      ≜ (∃v ∈ Value : Snd(v)) ∨ Rcv
vars      ≜ ⟨ch, tmt, rcv⟩
Chan      ≜ Init ∧ □[Next]vars

```

Refined specification of producer/consumer (1/2)

MODULE *RefinedPC*

EXTENDS *Sequences*

CONSTANTS *Product, None, buffCap*

VARIABLES *buffer, in, itmt, ircv, out, otmt, orcv*

InChan \triangleq INSTANCE *Channel* WITH

Value \leftarrow *Product*, *NoValue* \leftarrow *None*, *ch* \leftarrow *in*, *tmt* \leftarrow *itmt*, *rcv* \leftarrow *ircv*

OutChan \triangleq INSTANCE *Channel* WITH

Value \leftarrow *Product*, *NoValue* \leftarrow *None*, *ch* \leftarrow *out*, *tmt* \leftarrow *otmt*, *rcv* \leftarrow *orcv*

Init \triangleq *buffer* = $\langle \rangle$ \wedge *InChan*!*Init* \wedge *OutChan*!*Init*

Produce(*v*) \triangleq *InChan*!*Snd*(*v*) \wedge UNCHANGED \langle *buffer*, *OutChan*!*vars* \rangle

Get \triangleq \wedge *Len*(*buffer*) < *buffCap*
 \wedge *InChan*!*Rcv* \wedge *buffer*' = *Append*(*buffer*, *in*)
 \wedge UNCHANGED *OutChan*!*vars*

Put \triangleq \wedge *Len*(*buffer*) > 0
 \wedge *OutChan*!*Snd*(*Head*(*buffer*)) \wedge *buffer*' = *Tail*(*buffer*)
 \wedge UNCHANGED *InChan*!*vars*

Consume \triangleq *OutChan*!*Rcv* \wedge UNCHANGED \langle *buffer*, *InChan*!*vars* \rangle

Refined specification of producer/consumer (2/2)

$$\begin{aligned} \text{Next} &\triangleq (\exists v \in \text{Product} : \text{Produce}(v)) \vee \text{Get} \vee \text{Put} \vee \text{Consume} \\ \text{vars} &\triangleq \langle \text{buffer}, \text{InChan!vars}, \text{OutChan!vars} \rangle \\ \text{Fairness} &\triangleq \text{WF}_{\text{vars}}(\text{Get}) \wedge \text{WF}_{\text{vars}}(\text{Put}) \wedge \text{WF}_{\text{vars}}(\text{Consume}) \\ \text{Spec} &\triangleq \text{Init} \wedge \Box[\text{Next}]_{\text{vars}} \wedge \text{Fairness} \end{aligned}$$

LOCAL INSTANCE *AbstractPC*

THEOREM $\text{Spec} \Rightarrow \text{AbstractPC!Spec}$

• Proof of refinement theorem

- ▶ safety part quite obvious
- ▶ in particular, new actions do not modify variable *buffer*
- ▶ fairness condition on new *Consume* action necessary for implementing high-level fairness constraint on *Put*

• TLC can again verify refinement for small instances

Decomposition: partition variables and actions

- Decompose system into three components

- ▶ identify variables accessed by each component
- ▶ assign responsibilities for performing actions to components

	Producer	Buffer	Consumer
writes	<i>in, itmt</i>	<i>ircv, buffer, out, otmt</i>	<i>orc</i>
reads	<i>ircv</i>	<i>in, itmt, orc</i>	<i>out, otmt</i>
performs	<i>Produce</i>	<i>Get, Put</i>	<i>Consume</i>

- Derive component specifications from this decomposition

- ▶ no variable written by two different components
- ▶ unique assignment of actions to components

Producer and Consumer components in TLA⁺

MODULE *Producer*

CONSTANTS *Product, None*

VARIABLES *in, itmt, ircv*

PInit $\triangleq in = None \wedge itmt = \text{FALSE}$

Produce(*v*) $\triangleq \wedge itmt = ircv$
 $\wedge in' = v \wedge itmt' = \neg itmt \wedge ircv' = ircv$

Producer $\triangleq PInit \wedge \square [\exists v \in Product : Produce(v)]_{\langle in, itmt \rangle}$

MODULE *Consumer*

VARIABLES *orcvc, out, otmt*

CInit $\triangleq orcvc = \text{FALSE}$

Consume $\triangleq \wedge otmt \neq orcvc$
 $\wedge orcvc' = \neg orcvc \wedge \text{UNCHANGED } \langle out, otmt \rangle$

Consumer $\triangleq CInit \wedge \square [Consume]_{orcvc} \wedge \text{WF}_{orcvc}(Consume)$

Buffer component in TLA⁺

MODULE *Buffer*

EXTENDS *Sequences*

CONSTANTS *Product*, *None*, *buffCap*

VARIABLES *in*, *itmt*, *ircv*, *buffer*, *out*, *otmt*, *orc*

BInit \triangleq $ircv = \text{FALSE} \wedge buffer = \langle \rangle \wedge out = \text{None} \wedge otmt = \text{FALSE}$

Get \triangleq $\wedge Len(buffer) < buffCap \wedge itmt \neq ircv$
 $\wedge buffer' = \text{Append}(buffer, in) \wedge ircv' = \neg ircv$
 $\wedge \text{UNCHANGED } \langle in, itmt, out, otmt, orc \rangle$

Put \triangleq $\wedge Len(buffer) > 0 \wedge otmt = orc$
 $\wedge out' = \text{Head}(buffer) \wedge otmt' = \neg otmt \wedge buffer' = \text{Tail}(buffer)$
 $\wedge \text{UNCHANGED } \langle in, itmt, ircv, orc \rangle$

bvars \triangleq $\langle ircv, buffer, out, otmt \rangle$

Buffer \triangleq $BInit \wedge \Box [Get \vee Put]_{bvars} \wedge \text{WF}_{bvars}(Get) \wedge \text{WF}_{bvars}(Put)$

Summary: decomposition

MODULE *Refinement*

EXTENDS *Producer, Consumer, Buffer*

LOCAL INSTANCE *RefinedPC*

THEOREM $Producer \wedge Buffer \wedge Consumer \Rightarrow RefinedPC!Spec$

- Overall approach to decompose system specification
 - ▶ partition system variables according to who modifies them
 - ▶ identify actions performed by each component
 - ▶ declare all variables read or written in component specifications
 - ▶ only written variables appear in subscripts (next-state, fairness)
 - ▶ add UNCHANGED clauses for read-only variables
- What if no suitable variable partitioning can be found?
 - ▶ try to split (auxiliary) variables to separate responsibilities
 - ▶ allow for shared actions in different system models

Summary: structuring models

- Refinement: successively add implementation detail
 - ▶ represented in TLA as (validity of) implication
 - ▶ stuttering invariance allows for additional low-level steps
- Hiding of internal state components: exhibit system interface
 - ▶ represented in TLA by existential quantification over state variables
 - ▶ non-standard definition ensures stuttering invariance
 - ▶ standard proof rules remain sound: refinement mappings
- Composing and decomposing specifications
 - ▶ composition represented as conjunction of specifications
 - ▶ stuttering invariance allows for local component steps
 - ▶ decomposition derives components from system specification

Thank you!